

Digital Forensics

Federico Conti

2024/25

Contents

Forensic Acquisition	4
HDD and SSD technologies	4
HDD	4
SSD	5
Interface standards and protocols	6
ATA specification	6
Disk encryption: Bitlocker	8
Toolset and examples	9
losetup and dd	10
A faulty disk	11
dc3dd	12
Image file formats	12
Guymager	12
FTK Imager	13
File Systems	14
VSFS (Very Simple File-System)	14
Example	16
The Sleuth Kit (TSK)	16
Example	17
Example	17
DOS (or MBR) partition tables	18
Example	19
Extended partitions	20
Example	21
Example	22
GPT - GUID PARTition Tables	23
Example	23
File System Analysis	23
Example	24
Example	25
Example	26
TSK metadata commands	26
Example	27
Carving	29
Example	30
FAT	31
Volume Organization	31
Files and Directories	34
Example	35
Example	36
NTFS	37
Volume Organization	38
EXT	43
Layout	43
Example	45
Example	46
Inode	47
Directories	48
Example	50

Extended	51
--------------------	----

Forensic Acquisition

Acquisition is the process of cloning or copying digital data evidence.

- forensically sound (integrity and non-repudiation)
 - the copies must be identical to the original
 - the procedures must be documented and implemented using known methods and technologies, so that they can be verified by the opposite party
- a critical step
 - proper handling of data ensures that all actions taken on it can be checked, repeated, and verified at any time
 - incomplete or incorrect handling of data has the potential to compromise the entire investigation

It is generally recommended to avoid conducting analysis on the original device (namely, best evidence).

Creating a forensic image is typically considered the most effective method for preserving digital evidence (One or more, usually two).


Accessing the original media only once during the acquisition phase can help minimize the risk of altering or damaging the evidence

HDD and SSD technologies

HDD

A **hard disk** is a sealed unit containing a number of **platters** in a stack:

- each platter has two wothey moverking **surfaces**
- each working surface is divided into a number of concentric rings called **tracks**
 - the collection of all tracks that are the same distance from the edge of the platter is called a **cylinder**.
- each track is further divided into **sectors**
 - a sector is the smallest unit that can be accessed on a storage device
 - traditionally containing 512 bytes of data each, but recent hard drives have switched to 4KB sectors (Advanced Format)
 - a **cluster** is a group of sectors (from 1 to 64 sectors) that make up the smallest unit of disk allocation for a file within a file system



The data on a hard drive is read by read-write **hands**. The standard configuration uses one head per surface and they moves simultaneously from one cylinder to another.

A **low level format** is performed on the blank platters to create data structures for tracks and sectors

- creates all of the headers and trailers marking the beginning and ends of each sector
- header and trailer also keep the linear sector numbers (cf. LBA later), and error-correcting codes (ECC)

All disks are shipped with a few bad sectors (additional ones can be expected to go bad slowly over time).

- disks keep spare sectors to replace bad ones
- ECC calculation is performed with every disk read or write: if an error is detected but the data is recoverable, then a *soft error* has occurred if the data on a bad sector cannot be recovered, then a *hard error* has occurred. A bad sector can be replace with a spare one (but any information written is usually lost)

Older hard drives used a system called **CHS (Cylinder-Head-Sector)** addressing to locate data on the disk. This method relied on:

- Cylinders (C) → The track number (a ring of data on a disk platter).
- Heads (H) → The read/write head that accesses a platter's surface.
- Sectors (S) → The smallest unit of storage on a track.

Example: CHS (100, 2, 30) would mean: Cylinder 100; Head 2 (indicating the second platter side); Sector 30 (the exact location on that track).

CHS had physical limitations → It could only handle disks up to 504 MB (later ECHS extended this to 8 GB).

Instead of using three values (CHS), **Logical Block Addressing (LBA)** assigns a single number to each sector. LBA starts at 0 and counts up sequentially. The operating system and file system treat the disk as a continuous array of sectors, making it easier to manage.

Example: LBA 123456: The 123,456th sector on the disk.

SSD

A **Solid-State Drive (SSD)** is a non-volatile storage device, meaning it retains data even when the power is off. Unlike Hard Disk Drives (HDDs), which use spinning magnetic platters, SSDs rely on flash memory chips to store data.

- The smallest unit of an SSD is a **page**, which is composed of several memory cells: the page sizes are 2KB, 4KB, 8KB, 16KB or larger (usually they are 4 KB in size).
- Several pages on the SSD are summarized to a **block**: the block size typically varies between 256KB (128 pages * 2KB per page) and 4MB (256 pages * 16KB per page)

Operation	Description
Read and Write	An SSD can read and write data at the page level.
Write	Writing is only possible if other pages in the block are empty. Otherwise, it leads to write amplification.
Erase Data	An SSD can only erase an entire block at once due to the physical and electrical characteristics of the memory cells.

Modifying data requires a **Program/Erase (P/E) cycle**.

- During a P/E cycle, an entire block containing the targeted pages is written to memory
- The block is then marked for deletion, and the updated data is rewritten to another block

The erase operation does not happen immediately after data is marked for deletion. Instead, the SSD performs it asynchronously when necessary to optimize performance.

Garbage Collection helps free up space efficiently while minimizing interruptions to read/write operations.

Every flash memory cell has a limited number of P/E cycles before it wears out **Wear leveling** is a technique used by SSDs to distribute writes evenly across all memory blocks.

Unlike HDDs, SSDs cannot simply overwrite deleted files. If deleted data is not managed properly, unnecessary data copies can cause write amplification, increasing wear on the SSD. TRIM command allows the OS to inform the SSD which pages are no longer needed, enabling it to manage space more efficiently.

Forensic Investigators Face a Big Problem with SSDs

TRIM permanently deletes data by triggering the garbage collector, making data recovery impossible. The garbage collector operates independently within the SSD controller, meaning:

- Even a hardware write blocker (a forensic tool to prevent data changes) cannot stop it.
- Data can change midway or between acquisitions, complicating forensic investig

Interface standards and protocols

Disks are accessed using standard interfaces that define how they physically and logically connect to a computer system. Each standard improves data transfer speeds and fixes limitations from older technologies. A disk interface consists of two key components:

- Physical Connection → Defines the type of cable and connector used to attach the disk to the system.
- Logical Connection → Defines the protocol that controls how data is transferred between the disk and the computer.

ATA (Advanced Technology Attachment):

a widely used interface standard with multiple versions:

- ATA-1, ATA-2, ..., ATA-8 → Each new version improves speed and capacity.
- ATAPI (ATA Packet Interface) → Allows removable media (CD/DVD drives) to be connected using ATA but still uses SCSI commands for data transfer.

Type	Description
PATA (Parallel ATA)	Older, also known as IDE (Integrated Drive Electronics). Uses ribbon cables with multiple pins.
SATA (Serial ATA)	Modern standard introduced in ATA/ATAPI-7. Uses thin serial cables, improving speed and efficiency.
SATA 3.0	Supports speeds up to 6 Gbit/s (SATA-600). Common in modern HDDs and SSDs.
ATA-8	Introduced optimizations for SSDs, including TRIM support.

SCSI (Small Computer System Interface):

now replaced by SAS (Serial Attached SCSI) that is based on the SCSI standard, but uses a serial interface to connect storage. Better scalable and faster (supports data transfer rates of up to 24 Gbit/s).

NVMe (Non-Volatile Memory Express):

Uses a PCIe interface to connect storage devices to a computer. Commonly used for high-performance solid-state drives since it supports data transfer rates of up to 32 GB/s.

USB (Universal Serial Bus):

- uses a serial interface to connect storage devices to a computer
- mass storage is the standard protocol used for storage devices
- commonly used for external hard drives, flash drives, and other portable storage devices
- USB 3.1 Gen 1 standard supports speeds up to 5 Gbit/s, while the USB 3.1 Gen 2 standard, supports speeds up to 20 Gbit/s.

ATA specification

Introduced in ATA-3, **hard disk passwords** are an optional security feature designed to restrict unauthorized access to a hard drive. However, it is a lock mechanism, not encryption, meaning data remains unencrypted and could be accessed by other means.

There are two passwords:

1. User Password → Set by the owner.
2. **Master Password** → was designed so an administrator can gain access in case the user password was lost (every hard disk is initially supplied with an undocumented master password).

If passwords are being used, there are two modes that the disk can operate:

1. high security mode: the user and master password can unlock the disk
2. maximum-security mode: the user password can unlock the disk but the master password can unlock the disk after the disk content have been wiped

A protected HD will require the SECURITY_UNLOCK command to be executed with the correct password before any other ATA command. After the password has been entered, the disk works normally until the disk is powered on

Some ATA commands are still enabled on the HD when it is locked (so it may show as a valid disk when it is connected to a computer); however, trying to read data from a locked disk will produce an error

Setting the Password:


- Can be configured via BIOS settings.
- Linux users can manage HDD passwords using tools like hdparm.

Introduced in ATA-4, the **Host Protected Area (HPA)** is a hidden storage section on a hard disk that is not accessible to the operating system. It is used mainly by hardware vendors for system recovery files, diagnostics, or factory reset tools. A HPA is at the end of the disk and when used, it can be accessed by reconfiguring the hard disk.

Two ATA commands that return maximum physical addressable sectors

- READ_NATIVE_MAX_ADDRESS: return the maximum physical address
- IDENTIFY_DEVICE: return only the number of sectors that a user can access

To create an HPA, the SET_MAX_ADDRESS command is used to set the maximum address to which the user should have access (to remove it, use SET_MAX_ADDRESS = READ_NATIVE_MAX_ADDRESS).



the SET_MAX_ADDRESS command support different settings, e.g.,

- volatility bit: the HPA exist after the hard disk is reset or power cycled (otherwise the effect is permanent)
- locking command: prevents modification to the maximum address until next reset

when the BIOS requires to read/write some data in the HPA it uses SET_MAX_ADDRESS with volatility bit and locking.

It is possible to protect settings with a password (different from HD passwords).

The **Device Configuration Overlay (DCO)** was introduced in ATA-6 and allows manufacturers to limit the apparent capabilities of a hard disk. This feature enables backward compatibility with older systems but can also be exploited to hide data.

A computer uses the IDENTIFY_DEVICE command to check an HDD's specifications (size, features, supported commands). If a DCO is applied, the IDENTIFY_DEVICE command will not show the actual full disk size or features.

This is achieved using two special ATA commands:

- DEVICE_CONFIGURATION_SET → Modifies the DCO settings to restrict visible disk space or disable features.
- DEVICE_CONFIGURATION_RESET → Restores the original settings, removing the DCO restrictions.


*Example:

A 2 TB hard drive can be made to appear as a 500 GB drive, hiding the remaining 1.5 TB from the operating system.*

The DCO and HPA can co-exist on the same HDD (but DCO must be set first).

The `DEVICE_CONFIGURATION_IDENTIFY` command return the actual features and size of a disk:

- we can detect DCO if `DEVICE_CONFIGURATION_IDENTIFY` \neq `IDENTIFY_DEVICE`



At least three different methods for detecting HPA on Linux: `dmesg`, `hdparm`, and `disk_stat` (https://wiki.sleuthkit.org/index.php?title=Disk_stat)

Disk encryption: Bitlocker

BitLocker is a full-disk encryption feature in Windows that protects data using a multi-layered encryption system.

It makes use of symmetric encryption (by default, AES-128).

On modern systems, it is coupled with a Trusted Platform Module (TPM):

- the main functions of TPM are the generation, storage and secure management of cryptographic keys
- on a computer without TPM a password can be used (then BitLocker encryption will be just as secure as the password you set)

BitLocker uses different symmetric key:

1. raw data is encrypted with the **Full Volume Encryption Key (FVEK)**
2. FVEK is then encrypted with the **Volume Master Key (VMK)**
3. VMK is in turn encrypted by one of several possible methods depending on the chosen authentication type (that is, **key protectors** or TPM) and recovery scenarios

The use of intermediate key (VMK between FVEK and any key protectors) allows changing the keys without the need to re-encrypt the raw data in a case a given key protector is compromised or changed.


- When changing a key protector, a new VMK will be created and used to encrypt the old FVEK with the new VMK


BitLocker supports multiple key protector options, depending on the security needs and device type.

TPM only

- The TPM module (a hardware security chip) decrypts the VMK using a Storage Root Key (SRK) stored in the TPM.
- The SRK is only released if Secure Boot passes, ensuring the device boots with its original OS and configuration.
- The BitLocker volume is unlocked automatically during boot, before the user logs in.

View article.





TPM + PIN

- The TPM module will only release the VMK if the user enters a correct PIN during the pre-boot phase.
- If too many incorrect PIN attempts occur, the TPM will lock access to the encryption key, preventing brute-force attacks.

Key Takeaways:

- BitLocker is excellent for protecting against physical threats like device theft or unauthorized hard drive access.
- It does NOT protect against malware, ransomware, or unauthorized logins by users on the same computer.
- TPM + PIN is the most secure option to prevent unauthorized access, even if a device is stolen.

BitLocker poses a problem for forensic investigators, as all information on the drive will be encrypted, and therefore unreadable. Some methods for breaking BitLocker password are:

- the RAM dump/hibernation file/page file attack: this attack is universal, and works regardless of the type of protector. It dumps from the computer's volatile memory (and possibly in the page/hibernation file) the VMK that is loaded unencrypted while the volume is mounted
- BitLocker recovery keys: in many situations recovery keys are stored in the user's Microsoft Account. Extracting those keys from their account allows instantly mounting or decrypting protected volumes regardless of the type of protector

Toolset and examples

A **loop** device is a special kind of block device that does not map to a physical hardware device (such as a hard disk) but instead maps to a regular file stored within a filesystem.

- useful to access a forensic image
- read/only can be forced
- the offset parameter could be useful to directly access a volume
- (can be used to simulate a block device to be acquired)

Key losetup Commands:

- `losetup -a` → Shows the status of all loop devices.
- `losetup -d [device]` → Detaches a loop device.
- `losetup -f` → Finds the first available (unused) loop device.
- `losetup -o [offset]` → Starts reading data at a specific offset in the file.
- `losetup -r /dev/loop0 [srcfile]` → Sets up a read-only loop device.

dd is the precursor of all acquisition tools, allowing for the acquisition of data bit by bit in raw format.

Key dd Options:

- if= → Input file (or device to copy from).
- of= → Output file (or device to write to).
- bs= → Block size (how much data to read/write at a time).
- conv= → Specifies conversion options (e.g., noerror to continue on errors).

Example

```
dd if=/dev/sda of=/mnt/dest/image.dd bs=512
```

losetup and dd

In this exercise, we will simulate a block device using a compressed forensic image and interact with it as if it were a real disk.

```
wget https://github.com/enricorusso/DF_Exs/raw/main/acquisition/image.dd.gz
gunzip image.dd.gz
```

```
# Before setting up the loop device, find an available one using
losetup -f
```

```
sudo losetup -r /dev/loop1 image.dd # Now, set up the image as a read-only loop device
```

```
sudo dmesg | tail # To verify that the loop device was correctly attached, check system logs
[84058.342422] loop1: detected capacity change from 0 to 2033664
```

```
sudo fdisk -l /dev/loop1 # To inspect the loop device and view partition details
```

```
# Since the loop device represents an entire disk,
# Linux does not automatically recognize partitions.
# Use partx to make them available
sudo partx -a /dev/loop1
```

```
mkdir -p /mnt/forensic_image
sudo mount -o ro /dev/loop1p1 /mnt/forensic_image
```

```
ls -l /mnt/forensic_image
```

```
#clean
sudo umount /mnt/forensic_image
sudo losetup -d /dev/loop1
```

In digital forensics, verifying the integrity of a forensic image is crucial to ensure that the data remains unchanged during analysis. This is done by calculating cryptographic hash values (MD5 and SHA1) before and after mounting the image → Before using the forensic image, compute its MD5 and SHA1 hashes.

```
md5sum image.dd
446144a4af914d7e55603b6042f20db1  image.dd

sha1sum image.dd
99540f5aaa170afb722729e980fd6dc34ff323
image.dd

md5sum /dev/loop1
446144a4af914d7e55603b6042f20db1
/dev/loop1

sha1sum /dev/loop1
99540f5aaa170afb722729e980fd6dc34ff323
/dev/loop1
```

The dd tool does not calculate hashes during acquisition, so forensic best practices require manually computing hashes before and after imaging to ensure data integrity.

```
# Instead of separately computing hashes before and after, we can stream data from dd to tee, simultane
sudo dd if=/dev/loop1 bs=512 | tee image.dd
| hashdeep -c md5,sha1 > image.src_hash #apt install hashdeep

# dd if=/dev/loop1 bs=512 → Reads data from the loop device.
# tee image.dd → Writes data to image.dd while also passing it to the next command.
# hashdeep -c md5,sha1 → Computes MD5 and SHA1 hashes as data is written.
# > image.src_hash → Saves the computed hashes to image.src_hash.

#Finale verification hash
md5sum image.dd
```

A faulty disk

In this exercise, we simulate a faulty disk by mapping a logical block device and introducing bad sectors. We then attempt to acquire it using dd, handling errors properly to maintain forensic integrity.

1. We create a logical “faulty” device (1Kb) with the command dmsetup*
 - 8 8 error → [starting sector; add sector] maps the next 8 sectors of 512 byte (8 to 16) of the bad_disk device to an error area. This means that any attempt to read or write to bad_disk sectors 8 to 16 will generate an error.
 - /dev/loop1 is the origin and must be initialized with a .dd (sudo losetup /dev/loop0 image.dd/)

```
sudo dmsetup create bad_disk << EOF
0 8 linear /dev/loop0 0
8 8 error
16 2033648 linear /dev/loop0 16
EOF
```

2. Scan the simulated bad sectors.

```
sh sudo badblocks -b 512 -v /dev/mapper/bad_disk
```

3. To ensure that all reads go directly to the faulty device (and are not cached), we disable readahead. Then, check the block device size:

```
sudo blockdev --setra 0 /dev/mapper/bad_disk
sudo blockdev --getsz /dev/mapper/bad_disk
```

4. Now, try acquiring the faulty disk with dd

```
sudo dd if=/dev/mapper/bad_disk of=bad.dd bs=512
```

Problem: dd stops when it hits a bad sector, preventing a complete acquisition.

5. To log bad sectors and replace them with zeros, use conv=sync,noerror

- tee bad.dd → Writes the output to bad.dd while streaming it to hashdeep for hashing.

```
sudo dd if=/dev/mapper/bad_disk bs=512 conv=sync,noerror
| tee bad.dd | hashdeep -c md5,sha1 > bad_image.src_hash
```

6. After acquisition, compare the hash of bad.dd to the hash calculated during acquisition in bad_image.src_hash

```
hashdeep -c md5,sha1 bad.dd
```

dc3dd

An enhanced version of dd designed specifically for digital forensics. It was developed by the DoD Cyber Crime Center (DC3) and includes several critical forensic features missing in standard dd.

Example

```
sudo dc3dd if=/dev/mapper/bad_disk of=bad.dd ssz=512 log=image.log hlog=hash.log hash=md5 hash=sha1
```

Image file formats

When acquiring digital evidence, the choice of image format is crucial for integrity, compatibility, and efficiency in analysis. There are two main categories of forensic image formats.

The output from dd acquisition is a raw image * it contains only the data from the source device * all the descriptive data about the acquisition (e.g., hashes values, dates, or times) need to be saved in a separate file

An embedded image contains data from the source device and additional descriptive data (metadata).

- **Expert Witness Format (EWF)**

Joachim Metz (Google) created the libewf project, open source (<https://github.com/libyal/libewf>, apt install ewf-tools). It provides a library and set of tools to manage the ewf format.

- ewfacquire: reads storage media data from devices and write files to EWF files.
- ewfexport: exports storage media data in EWF files to (split) RAW format or a specific version of EWF files.
- ewfinfo: shows the metadata in EWF files.
- ewfmount: FUSE mounts EWF files.
- ewfrecover: special variant of ewfexport to create a new set of EWF files from a corrupt set.
- ewfverify: verifies the storage media data in EWF files

(FUSE (Filesystem in Userspace) is a Linux kernel module that allows users to mount and manage file systems without requiring root privileges or kernel modifications.) → alternative :uses `sudo ewfmount`

- **Advanced Forensic Format (AFF)**

Open Source format developed by Dr. Simson L. Garfinkel

- Provide compressed or uncompressed image files
- No size restriction for disk-to-image files
- Provide space in the image file or segmented files for metadata (unlimited number)
- Digital signatures
- Encryption with decryption on-the-fly
- No patents

Still lacks wide adoption (software available at <https://github.com/sshock/AFFLIBv3>).

Guymager

Guymager is a graphical (Qt-based) forensic imager. It is capable of producing image files in EWF, AFF and dd format (apt install guymager).

- AFF is disabled by default (`sudo nano /etc/guymager/guymager.cfg` → set `AffEnabled=true` → `sudo systemctl restart guymager`)

FTK Imager

FTK Imager is a data preview and imaging tool used to acquire data (evidence) in a forensically sound manner by creating copies of data without making changes to the original evidence (<https://www.exterro.com/ftk-product-downloads>).

File Systems

In a computer, data storage is organized in a hierarchical manner.

At the top, we have fast and directly accessible storage, such as registers and cache, while at the bottom, we have slower but larger storage, like hard drives or SSDs.

We focus on secondary (external) memory storage:

- not directly accessible by CPU → data must be transferred to main memory (RAM) before it can be processed
- data transferred in blocks → rather than individual bytes
- significantly slower
- non-volatile → retains data even when power is turned off

Floppies/HDs/CDs/DVDs/BDs/SD-cards/SSDs/pendrives. . . are all block devices; following File System Forensic Analysis's terminology.

A **volume** is a collection of addressable blocks

- these blocks do not need to be contiguous on a physical device
- a volume can be assembled by merging smaller volumes

A **partition** is a contiguous part of a volume


- partitioning is optional: some removable storage does not use it

By definition, both disks and partitions are volumes. In this part of the course we deal with block-device (forensics) images, like the one acquired from actual devices.

Users don't interact directly with storage blocks—instead, they work with files and directories.

The **file system** creates this illusion; i.e., it handles the mapping between files/directories and a collection of blocks (usually, clusters of sectors). Consists of on-disk data structures to organize both data and metadata. There exist various file-system formats (e.g., FAT, NTFS, . . .)

(high-level) **formatting** a volume means to initialize those structures.




VSFS (Very Simple File-System)

In Unix-like file systems (e.g., EXT4, see `man mkfs.ext4`), each file (or other filesystem object) has an associated **inode** that stores its metadata. However, inode does not store file names, which are instead kept in directory structures.

Every file system object has an inode, including:


- regular file
- directory
- symbolic link
- FIFO
- socket
- character device
- block device

Formatting means preparing: the superblock, i-node/data bitmaps, i-node table, data region.



An inode contains:


- file type
- UID/GID/permission-bits
- time information
- size in bytes
- number of hard links (AKA names)
- pointers to data blocks



To use the end-user view, a file system, stored on a **block device**, must be **mounted** (or “parsed”)

Most modern operating systems automatically mount external storage devices when they are connected.

- In Unix/Linux there is a single root directory (/), and additional volumes are mounted within this hierarchy
- in Windows each volume (storage device/partition) is assigned a drive letter (C:, D:, E:)



Block devices can be seen as “files” themselves

- Linux special files, typically under /dev
 - various “aliases” in /dev/disk -> /by-id; /by-uuid; /bypath
 - lsblk lists information about available block devices

Viceversa, (image) files can be seen as block devices.

1. losetup command allows an image file to be treated as a virtual block device (loop device).

- --list
- --find [--show] [--partscan] image
- --detach[-all]

2. Then, we can mount them.

- Instead of manually setting up a loop device, mount can automatically create one:
- offset=<byte_offset> starting point within an image file (use fdisk -l image_file.img)
- mount [-o loop] image instead of manually setting up a loop device
- ro read-only

3. Umount and check umount /dev/sda1 fsck /dev/sda1

Example

```
xz -dk two-partitions.dd.xz
```

```
# FIRST METHOD
```

```
losetup -r -o $((1*512)) /dev/loop0 two-partitions.dd # First partition
losetup -r -o $((1026*512)) /dev/loop1 two-partitions.dd # Second partition
```

```
# losetup -r -o $((1*512)) --find --show /tmp/two-partitions.dd
# losetup -r -o $((1026*512)) --find --show /tmp/two-partitions.dd
```

```
fdisk -l /dev/loop1 # check
```

```
mount -o ro /dev/loop0 /mnt/two-partition
mount -o ro /dev/loop1 /mnt/two-partition
```

```
# SECOND METHOD
```

```
losetup -r --find --show --partscan two-partitions.dd
```

```
mount -o ro /dev/loop0p1 /mnt/part1
mount -o ro /dev/loop0p2 /mnt/part2
```

```
umount /dev/loop0p1
umount /dev/loop0p1
```

The Sleuth Kit (TSK)

!attention: TSK always uses a term 'inode', but actually the information of filesystem is in the FAT entry directory

The Sleuth Kit (TSK) is a forensic toolkit that provides different layers of analysis for digital investigations. Each layer focuses on specific aspects of a digital storage system, allowing forensic examiners to extract and interpret data at various levels.

- `img_` for images
- `mm` (media-management) for volumes
- `fs` for file-system structures
- `j` for file-system journals
- `blk` for blocks/data-units
- `i` for inodes, the file metadata
- `f` for file names

Typically followed by:

- `stat` for general information
- `ls` for listing the content
- `cat` for dumping/extracting the content

Example

```
img_stat two-partitions.dd
img_cat two-partitions.dd
```

```
img_stat canon-sd-card.e01
```

When analyzing file systems, we categorize data into essential and non-essential based on their reliability and importance.

- Essential Data = Trustworthy & required for file retrieval.
 - If name or location were incorrect, then the content could not be read
- Non-Essential Data = Can be misleading & needs verification.
 - the last-access time or the data of a deleted file could be correct but we don't know

The Volume (or Media Management) layer in The Sleuth Kit (TSK) focuses on analyzing and managing disk partitions. This layer is crucial for identifying partition structures, extracting partitions, and verifying file system integrity.

- `mmstat image` displays the type of partition scheme
- `mmls image` displays the partition layout of a volume
- `mmcat image part_num` outputs the contents of a partition

Example

For canon-sd-card.e01

1. Find the type of partition table (`mmstat`)
2. List the partitions (`mmls`)
3. Extract the DOS FAT16 partition, by using both `mmcat/dd` or a dd-like tool

Check whether the SHA256 of their results match Read-only mount the FAT partition and list the files

```
mmstat canon-sd-card.e01
mmls canon-sd-card.e01
```

```
##OUT##
```

```
DOS Partition Table
```

```
Offset Sector: 0
```

```
Units are in 512-byte sectors
```

	Slot	Start	End	Length	Description
000:	Meta	0000000000	0000000000	0000000001	Primary Table (#0)
001:	-----	0000000000	0000000050	0000000051	Unallocated
002:	000:000	0000000051	0000060799	0000060749	DOS FAT16 (0x04)

```
###
```

```
# First method (TSK toolkit)
```

```
mmcat canon-sd-card.e01 2 > fat16_mmcat.e01
```

```
# Second method
```

```
ewfmount canon-sd-card.e01 ./rawimage/ # bit a bit copy
```

```
sudo dd if=rawimage/ewf1 of=fat16_dd.dd bs=512 skip=51
```

```
sudo umount rawimage
```

```
sha256sum fat16_mmcat.dd fat16_dd.dd # equals
```

```
# First method (TSK toolkit)
```

```
fls -r -o 51 canon-sd-card.e01 #
```

```
#Second method
```

```
mount -o ro fat16_dd.dd /mnt/fat16_dd
```


```
tree /mnt/fat16_dd
```

DOS (or MBR) partition tables

The concept of MBR was introduced in 1983 with PC DOS 2.0.

It contain:

- machine code for the [boot loader](#), which usually loads and executes the active-partition [Volume Boot Record](#)
- a 32-bit unique identifier for the disk, located at offset 440 (0x1B8).
- information on how the disk is partitioned → four 16-byte entries (each at offset 446 (0x1BE)), allowing up to four primary partitions.
- last two bytes of the MBR contain the signature bytes: 0x55 0xAA.



1. Valid Configurations (A, B, and C):

- These configurations ensure that partitions are either adjacent or properly aligned without overlap.
- Partitions are defined in a way that does not create ambiguity in data storage.

2. Invalid Configurations (D and E):

- D and E depict overlapping partitions, which is problematic.
- Overlapping partitions may cause data corruption, boot issues, or system conflicts because two partitions would claim the same disk space.

[CHS \(Cylinder-Head-Sector\)](#) is the early method for addressing physical blocks on a disk.

It used a 3-byte structure:

- 10 bits for Cylinders (tracks stacked vertically)
- 8 bits for Heads (read/write heads on a disk platter)
- 6 bits for Sectors (sections of a track)

Replaced by [Logical Block Addressing](#) in '90s.

- To convert you need to know the number of heads per cylinder, and sectors per track, as reported by the disk drive
- Yet, many tools still aligned partitions to cylinder boundaries

	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F
0	Master Boot Code															
...																
1A0																
1B0											Disk Signature				Boot ind ¹	Start head
1C0	start Sect ²	start Cyl ²	Sys ID ³	End Head	End sect ²	End Cyl ²	Relative Sectors				Total Sectors					
1D0																
1E0																
1F0															55	AA

1. Boot indicator 0x00 = non-boot, 0x80 = bootable

2. Starting sector & starting cylinder are allocated bits, not bytes (0x1C0-0x1C1) same goes for end head and end sector

BIT	0	1	2	3	4	5	6	7	8	9	A		B	C	D	E	F
Value	Starting sector						Starting Cylinder										

3. Common partition values.

0x01	FAT12 <32MB
0x04	FAT16 <32MB
0x05	MS Extended partition using CHS
0x06	FAT16B
0x07	NTFS, HPFS, exFAT
0x0B	FAT32 CHS
0x0C	FAT32 LBA
0x0E	FAT16 LBA
0x0F	MS Extended partition LBA
0x42	Windows Dynamic volume
0x82	Linux swap
0x83	Linux

0x84	Windows hibernation partition
0x85	Linux extended
0x8E	Linux LVM
0xA5	FreeBSD slice
0xA6	OpenBSD slice
0xAB	Mac OS X boot
0xAF	HFS, HFS+
0xEE	MS GPT
0xEF	Intel EFI
0xFB	VMware VMFS
0xFC	VMware swap

A Master Boot Record (MBR) is typically 512 bytes and laid out like this:

Offset (hex) | Size | Description

0x000	446	Bootstrap code area
0x1B8	4	Disk signature (sometimes called "unique MBR signature")
0x1BC	2	Usually 0x0000 or may be used for copy-protection, etc.
0x1BE	16	Partition entry #1
0x1CE	16	Partition entry #2
0x1DE	16	Partition entry #3
0x1EE	16	Partition entry #4
0x1FE	2	MBR signature (0x55AA)

Each 16-byte partition entry has the structure:

Byte | Description

0	Boot indicator (0x80 = bootable; 0x00 = non-bootable)
1-3	Starting CHS (Head-Sector-Cylinder) - often unused in modern disks
4	Partition type (ID)
5-7	Ending CHS (Head-Sector-Cylinder)
8-11	Relative sectors (start in LBA)
12-15	Total sectors in this partition

Example

Use ImHex, writing proper patterns, to extract disk and partition information from mbr{1,2,3}.dd. Then, answer the following questions:

1. What are the three disk signatures?

2. Is there any MBR with inconsistent partitioning?
3. Are there MBRs without bootable partitions?
4. What is the largest FAT (id=4) partition?
5. Are CHS information always present?

```
(fdisk -l mbr2.dd)
```

```
tar -tJf MBR123_and_GPT.tar.xz
tar -xJvf MBR123_and_GPT.tar.xz mbr1.dd
```

```
xxd -s 0x1B8 -l 4 mbr1.dd
```

Pattern editor

```
// fdisk give a same informations
```


```
#include <std/mem.pat>
struct PartitionEntry {
    u8  bootIndicator;
    u8  startCHS[3];
    u8  partitionType;
    u8  endCHS[3];
    u32 relativeSectors;
    u32 totalSectors;
};

struct MBR {
    u8 bootCode[0x1B8];           // 446 bytes
    u32 diskSignature;           // offset 0x1B8
    u16 reserved;                // offset 0x1BC (often 0x0000)
    PartitionEntry partitions[4]; // 4 partition entries, each 16 bytes
    u16 signature;               // offset 0x1FE, should be 0x55AA
};

MBR seg[while(!std::mem::eof())] @ 0x00;
```

Extended partitions

MBR has only 4 slots for primary partitions.



To work around this limitation, one slot can be used for the [primary extended partition](#), a partition containing other partitions.

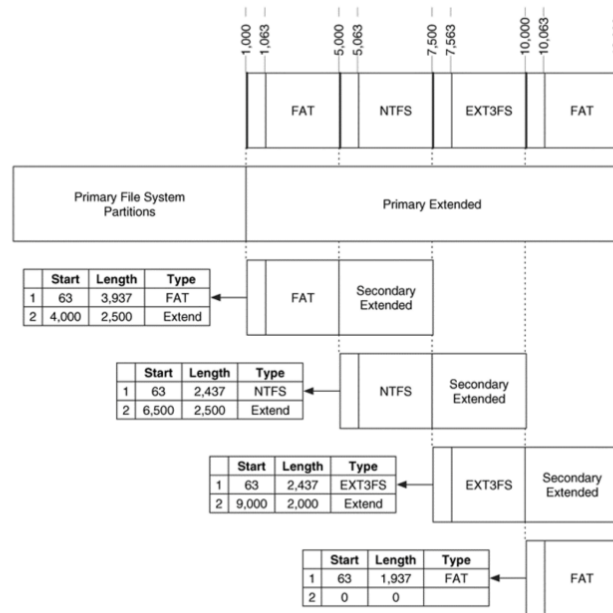
Beware of logical-partition addressing, which uses the distance from the beginning of a partition (vs the physical-addressing, from the beginning of the whole disk).

Inside the primary extended partition we find secondary extended partitions, containing

- a **partition table t** (with the same 512-byte structure)
- a **secondary file-system partition p** (logical partition), which contains a FS or other data

The partition table (t) describes:

1. Location of p (logical partition) relative to t.
2. Next **secondary extended partition** (if any), w.r.t. the primary extended partition



Example

ext-partitions.dd (SHA256: b075ed83211...) contains three partition tables: one primary, two extended. Analyze them with ImHex, and compare the result w.r.t. fdisk/mmls Source: (<https://dfit.sourceforge.net/test1/index.html>)

```
mmls ext-partitions.dd
```

	Slot	Start	End	Length	Description
000:	Meta	0000000000	0000000000	0000000001	Primary Table (#0)
001:	-----	0000000000	0000000062	0000000063	Unallocated
002:	000:000	0000000063	0000052415	0000052353	DOS FAT16 (0x04)
003:	000:001	0000052416	0000104831	0000052416	DOS FAT16 (0x04)
004:	000:002	0000104832	0000157247	0000052416	DOS FAT16 (0x04)
005:	Meta	0000157248	0000312479	0000155232	DOS Extended (0x05) #15724*512 = address
006:	Meta	0000157248	0000157248	0000000001	Extended Table (#1)
007:	-----	0000157248	0000157310	0000000063	Unallocated
008:	001:000	0000157311	0000209663	0000052353	DOS FAT16 (0x04)
009:	-----	0000209664	0000209726	0000000063	Unallocated
010:	001:001	0000209727	0000262079	0000052353	DOS FAT16 (0x04)
011:	Meta	0000262080	0000312479	0000050400	DOS Extended (0x05)
012:	Meta	0000262080	0000262080	0000000001	Extended Table (#2)
013:	-----	0000262080	0000262142	0000000063	Unallocated
014:	002:000	0000262143	0000312479	0000050337	DOS FAT16 (0x06)

Example

Someone purposely damaged the partition table of hidden-truth.dd (SHA256: 5f39a8965ec...)

1. Can you (ro) mount the partitions?
2. Can you repair the broken MBR and mount the deleted partition?

hint (ROT13): Lbh pna ernfba nobhg gur ynlbgh be trg fbzr uryc jvgu fvtsvaq (sebz GFX)

3. Can you recover the password protected "secret"?

```
fdisk -l hidden-truth.dd
```

```
##OUT##
Device          Boot      Start         End      Sectors  Size Id Type
hidden-truth.dd1             2         2050        2049     1M  4 FAT16 <3
hidden-truth.dd2    1751214177 2311246017 560031841   267G 74 unknown
hidden-truth.dd3          3076         8191        5116    2.5M  4 FAT16 <3
Partition table entries are not in disk order.
###
```

```
dd if=hidden-truth.dd of=hid2.dd bs=512 skip=3076 count=5115
dd if=hidden-truth.dd of=hid1.dd bs=512 skip=2 count=2048
```

```
ls -l hidden-truth.
```

```
##OUT##
-rwxrwxrwx 1 vagrant vagrant 4194304 Mar 13 19:37 hidden-truth.dd
# echo $((4194304/512)) == 8192 sectors
###
```

```
# try to find a sector 1 (0 is a boot sector)
dd if=hidden-truth.dd bs=512 skip=1 count=1 | xxd -g1
```

```
##OUT##
.....
000001e0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
000001f0: 34 37 36 39 36 66 37 36 33 34 36 65 36 65 32 31 47696f76346e6e21
# only clue ...
###
```

```
# try a "brute" mount
mount -o ro, offset=$((2051*512)) hidden-truth.dd /mnt/hidden-brute
```

```
##OUT##
not_so_secret.zip # a zip with pass
###
```

```
# does the fs cover the whole prtion in the middle?
# it is true that there is some space between the first and third partition but
# we do not know if all the space in between has been used
fsstat hidden-truth.dd -o 2051
```

##OUT##

File System Layout (in sectors)

Total Range: 0 - 1023 # 3075-2051 == 1024, okey seems fair

###

using a hex before and cyberchef

47696f76346e6e21 --> Giov4nn!

GPT - GUID PARTition Tables

A Universally/Globally Unique Identifier (UUID/GUID) is a 128-bit label.

- Uniqueness: Properly generated UUIDs are statistically unique, meaning the probability of duplication is extremely low.
- Standard Format: UUIDs are typically written in a 32-character hexadecimal format divided into five groups: 8-4-4-4-12, separated by hyphens.

`uuidgen`

bdeec955-b1b8-44a2-8034-15507d431aca

The GPT format, used by the Extensible Firmware Interface (EFI), which replaced BIOS, is the current standard on PCs; it

- starts with a protective MBR
- supports up to 128 partitions
- uses 64-bit LBA addresses
- keeps “mirrored” backup copies of
- important data structures

Example

Use ImHex, writing proper patterns, to extract disk and partition information from gpt.dd. Then, answer the following questions:

1. What is the disk GUID?
2. How many partitions are there?
3. What are the partition names?
4. Can you find the partition type GUIDs in the previous table?

```
gdisk -l gpt.dd
```

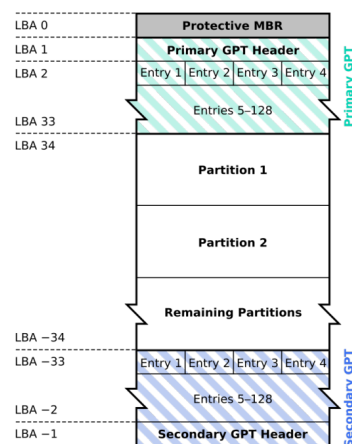
```
mmls -t gpt gpt.dd
```

File System Analysis


A reference model for a file system based on different categories of data that are involved in file storage and management.

- File System Category: layout and size information about the entire file system, such as: file system parameters (e.g., block size, total size) the structure or mapping of data storage
- Content The actual data, stored in clusters/blocks/data-units
- MetaData: Data that describes files: size, creation date

GUID Partition Table Scheme



- File Name Data that assign names to files
- Application: Data not needed for reading/writing a file; e.g., user quota statistics or a FS journal



To get the general details of a file-system - `fsstat [-o sect_offs] image`

Example

1. Find the OEM Name and Volume Label (Boot Sector) in `canon-sd-card.e01`

`mmls canon-sd-card.e01`

##OUT##

Slot	Start	End	Length	Description
000: Meta	0000000000	0000000000	0000000001	Primary Table (#0)
001: -----	0000000000	0000000050	0000000051	Unallocated
002: 000:000	0000000051	0000060799	0000060749	DOS FAT16 (0x04)

###

`fsstat -o 51 canon-sd-card.e01`

2. Check whether the partition types are correctly set inside `two-partitions.dd`

`mmls two-partitions.dd`

##OUT##

Slot	Start	End	Length	Description
000: Meta	0000000000	0000000000	0000000001	Primary Table (#0)
001: -----	0000000000	0000000000	0000000001	Unallocated
002: 000:000	0000000001	0000001025	0000001025	DOS FAT16 (0x06) # wrong
003: 000:001	0000001026	0000002047	0000001022	DOS FAT12 (0x01)

Partition table can be modified

###

`fsstat -o 1 two-partitions.dd # FAT12`

`fsstat -o 1026 two-partitions.dd # FAT12`

Example

inside the image file two-partitions.dd

1. look for the strings

- "didattica"
- "wDeek"
- "tool"
- "secret"

```
strings two-partitions.dd | grep -E "didattica|wDeek|tool|secret" # secret,wDeek,didattica
```

```
or
```

```
strings two-partitions.dd | ag "didattica|wDeek|tool|secret"
```

```
or
```

```
xxd -g1 two-partitions.dd | grep -C 3 ecre
```

2. (ro) mount its partitions, and look for the same strings inside the contained files

```
losetup -r --find --show --partscan two-partitions.dd
```

```
mount -o ro /dev/loop0p1 /mnt/part1
```

```
mount -o ro /dev/loop0p2 /mnt/part2
```

```
grep -rE "didattica|wDeek|tool|secret" /mnt/part1 # null
```



```
grep -rE "didattica|wDeek|tool|secret" /mnt/part2 # tool
```

Do some string appear only in one of the two searches? Can you guess why?

Each sector can have multiple addresses, relative to the start of the...

- storage media: physical address
- volume: (logical) volume address
- FS [data area]: (logical) FS address AKA (logical) cluster number
- file: (logical) file address AKA virtual cluster numbers

When writing a 612-byte file in a file system with 2K clusters (where each sector is 512 bytes), the way data is allocated creates slack space—unused but allocated storage that may contain remnants of previous data.



When investigating deleted files, forensic analysts use two major approaches:

1. Metadata-based

If the file is deleted but metadata still exists, we can recover:

- File size, timestamps, and allocated sectors/clusters.
- Orphaned files (files with no full path reference)

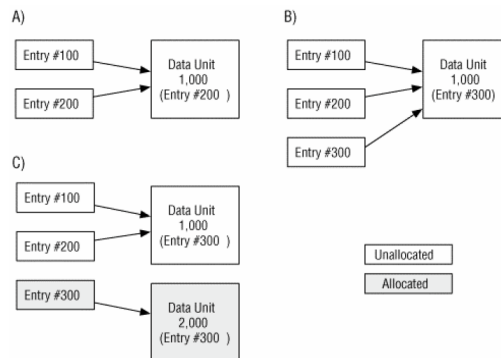
2. Application-Based

Used when metadata is unavailable:

- Typically from un-allocated space
- Does not need any FS information

Example

Where do data-units come from?



A.

- Entry#100 initially points to Data Unit 1,000.
- Entry #200 is created after #100 is deleted and reuses the same data unit.

This means that even after deletion, old data might still be recoverable unless overwritten.

C.

- Entry #300 is now assigned a completely new Data Unit (2,000).
- Entry #100 and #200 had used Data Unit 1,000, but it is now unallocated.

The original content in Data Unit 1,000 might still be present but no longer linked to any active file. (Carving).

TSK metadata commands

The Sleuth Kit (TSK) provides powerful commands to analyze file system metadata, particularly focusing on inodes, which store key file attributes.

1. `ils [-o sect_offs] image` - list inode information

- `-r` → Lists only removed (deleted) files
- `-a` → Lists only allocated (active) files
- `-m` → Displays inode details in a format compatible with mactime (used for timeline analysis)

2. `istat [-o sect_offs] image inum` - dumps detailed metadata of a specific file or inode

!!! TSK uses the inode abstraction even for file systems that do not natively have them.

- Some file systems (e.g., FAT32) do not have inodes, but TSK emulates them to allow a consistent analysis approach

3. `ifind [-n filename] [-d data-unit] [-o offset] image` - viceversa, to find the inode corresponding to a data-unit or file name

- `strings -t d disk-image.dd | grep "password"` - gives you an offset (e.g., 123456) where the data appears.
- `ifind -d $((123456/4096)) disk-image.dd - $((n/block-size))` returns the inode number

4. `ffind [-o sect_offs] image inum` - lists the names using the inode (useful when names are not inside the "inode")
5. `icat [-o sect_offs] image inum` - extracts and displays the contents of a file based on its inode number.

- `-s` → Includes slack space (unused space in the last cluster of a file)
- `-r` → Attempts to recover deleted files

Note: deleted content may be present in unallocated data (without metadata pointing to it). To check/dump blocks:

- `blkstat image block` - displays metadata about a specific block (e.g., allocation status, timestamps, etc.)
- `blkcat image block [how-many-blocks]` - outputs the raw content of a specific block
- `blkls` - lists or outputs blocks too

6. `fls [-o sect_offs] image [inum]` - list files inside the directory corresponding to the inode number
7. `ffind [-o sect_offs] image inum` - lists the names using the inode (useful when names are not inside the "inode")

Example

Let's find out why some of the following strings appear in one search and not the other

`mmls two-partitions.dd`

```
##OUT##
      Slot      Start      End      Length      Description
000:  Meta      0000000000  0000000000  0000000001  Primary Table (#0)
001:  -----      0000000000  0000000000  0000000001  Unallocated
002:  000:000      0000000001  0000001025  0000001025  DOS FAT16 (0x06)
003:  000:001      0000001026  0000002047  0000001022  DOS FAT12 (0x01)
###
```

`strings -t d two-partitions.dd | grep -E "didattica|wDeek|tool|secret"`

```
##OUT##
20514 and I have a secret message ;)
547396 wDeek
547436 /home/gio/didattica/file-systems/vol_fs_analysis/examples/pp/test
###
```

- End of first partition = $(1025 * 512) = 524800$ → "secret" is in the first partition.
- Start of second partition = $(1025 * 512) = 525312$ → "wDeek" is in the second partition.
- End of second partition = $(1025 * 512) = 1048064$ → "didattica" is in the second partition.

1. Find "secret"

- Sector number of "secret" = $(20514 / 512) = 40$ at beginning of the disk, but partition start at sector 1
- Offset = $(40 - 1) = 39$

`ifind -d 39 -o 1 two-partitions.dd # get 4 (a inode of block 39 of partition start 1)`

`istat -o 1 two-partitions.dd 4`

```
##OUT##
```

```
Directory Entry: 4
Allocated
File Attributes: File, Archive
Size: 34
Name: HELLO.TXT
```

```
Directory Entry Times:
Written:      2023-03-16 08:57:32 (EDT)
Accessed:     2023-03-16 00:00:00 (EDT)
Created:      2023-03-16 08:57:32 (EDT)
```

```
Sectors:
39 0 0 0 # use only one sector
###
```

```
icat -s -o 1 two-partitions.dd 4
```

```
##OUT##
Hi there! ...
###
```

We note that the size of ls hello is (34B) < oh the size dd rows (64)

```
dd if=two-partitions.dd bs=512 count=1 skip=40 | hexdump -C
ls -l hello.txt
```

2. Find “wDeek” and “didattica”

- Sector number of “secret” = $(547396/512) = 1069$ **at beginning of the disk, but partition start at sector 1026**
- Offset = $(1069-1026) = 43$

```
ifind -d 43 -o 1026 two-partitions.dd # get 6 (a indd of block 43 of partition start 1026)
```

```
istat -o 1026 two-partitions.dd 6
```

```
##OUT##
Directory Entry: 6
Not Allocated # DELETED --> not mounted by OS
File Attributes: File, Archive
Size: 4096
Name: _EST~1.SWP
```

```
Directory Entry Times:
Written:      2023-03-16 09:04:10 (EDT)
Accessed:     2023-03-16 00:00:00 (EDT)
Created:      2023-03-16 09:04:10 (EDT)
```

```
Sectors:
43 44 45 46 47 48 49 50
###
```

```
icat -o 1026 two-partitions.dd 6 | strings
```

```
##OUT##
b0VIM 8.2
root
wDeek
/home/gio/didattica/file-systems/vol_fs_analysis/examples/pp/test
3210
#!

###
```

3. Find “tool”

```
fls -rp two-partitions.dd -o 1026
```

```
##OUT##
r/r 4:  wikipedia.txt
r/r * 6:      .test.swp
r/r * 8:      test
v/v 16083:    $MBR
v/v 16084:    $FAT1
v/v 16085:    $FAT2
V/V 16086:    $OrphanFiles
###
```

```
istat -o 1026 two-partitions.dd 4
```

```
##OUT##
Directory Entry: 4
Allocated
File Attributes: File, Archive
Size: 3934
Name: WIKIPE~1.TXT

Directory Entry Times:
Written:      2023-03-16 09:04:24 (EDT)
Accessed:     2023-03-16 00:00:00 (EDT)
Created:      2023-03-16 09:04:24 (EDT)

Sectors:
39 40 41 42 55 56 57 58
# we see that the cluster is not consecutive and string "tool"
# is fragmented in "to...ol"
###
```

```
icat -o 1026 two-partitions.dd 4 | strings | hexdump -C | grep -C 6 tool
```

Carving

Is a method of recovering files without relying on metadata (like file names, paths, or inodes). It works by identifying file signatures (headers & footers) and extracting the data between them.(E.g., 0xFF 0xD8 and 0xFF 0xD9 for JPEG files).

Example

Inside `eighties.dd` (`cc121c3a037f904a4fa5ef51263df9fdb800d89af7330df22615802b81821f9d`) there is a FAT file system with some deleted content. In particular, there were files with the following SHA256 hashes:

- 4410aaee5ae15917c064f80a073ec75260482b7035fad58c85f1063d0b795733
- 1b756ad00ad842c3356c093583e2e4fab2540e15ca88750606f45f7efd1f4d26
- 592f47dfcbda344fc394987b6e02a65a35d4d849d35d2fc821e5be1889c645d
- 8a461036c70736eb4ca83e9062318c8293be2baad1c475c41c1945221559048e
- 0d176b77f6b81468eb7ba367d35bdcdb8fdcf63445c2cc83c5e27c5e0b4c1a14

Can you recover and identify them?

```
fls -rp eighties.dd or ils -r eighties.dd
```

```
##OUT##
```

```
# seven deleted (*)
```

```
r/r * 3:      -
r/r * 4:      -
r/r * 5:      _8.gif
r/r * 6:      _8.txt
v/v 523203:   $MBR #
v/v 523204:   $FAT1
v/v 523205:   $FAT2
V/V 523206:   $OrphanFiles
-/r * 517:    $OrphanFiles/_live.jpg
-/r * 518:    $OrphanFiles/_8k.jpg
-/r * 581:    $OrphanFiles/_monty.tzx
###
```

```
icat eighties.dd 6 | sha256sum # ok
icat eighties.dd 5 | sha256sum # icat fail sha
```

```
icat eighties.dd 5 | xxd -g1 # magic is gif, but there is some text
icat eighties.dd 5 > fake.gif # mmm...blurred ...
```

```
# wait ... two files cannot share the same sector
```

```
istat eighties.dd 5
```

```
##OUT##
```

```
Sectors:
```

```
108 109 110 111 112 0 0 0
```

```
###
```

```
istat eighties.dd 6
```

```
##OUT##
```

```
Size: 2426
```

```
Name: _8.gif
```

```
Sectors:
```

```
112 113 114 0
```

```
###
```

```

# Two possible cases:
# 1. the file may have been overwritten (nothing can be done)
# 2. the gif was framed,
    # first it was created a txt after gif ...
    # the gif used space that was previously empty in the txt file and continued to use other space

# This information is saved in FAT, but when a file is deleted, the cluster chain is lost.
# then we can recover the first cluster and thanks the length we can find other cluster

# cluster size is 4 sectors
fsstat eighties.dd
##OUT##
Cluster Size: 2048 # 2048/512=4
###

img_cat -s 108 -e 111 eighties.dd > fake2.gif
img_cat -s 116 -e 119 eighties.dd >> fake2.gif # okay noe its clear
sha256sum fake2.gif # but not yet, because here we have taken two clusters
# the file must be size: 2426, these bytes actually make up the .gif file
dd bs=1 count=2426 if=fake2.gif of=speriamo.gif
sha256sum speriamo.gif # OK 1b756...

```

FAT

The FAT (File Allocation Table) File System is one of the earliest and simplest file systems, first developed in 1977/1978. Over time, it evolved into three main versions: FAT12, FAT16 and FAT32.

The number indicates the # of bits used to identify clusters

- FAT12 can address $2^{12} = 4096$ clusters. Windows permits cluster sizes from 512 bytes to 8 KB, which limits FAT12 to 32 MB
- FAT16 can address $2^{16} = 65,536$ clusters
- FAT32 can address 2^{28} clusters (top 4 bits used for other purposes)

Actually, first 2 & last 16 are reserved: usable clusters are slightly less.

Uses the MSDOS 8.3 filename format – Only 8 characters for the name + 3-character file extension (e.g., FILE1234.TXT).

VFAT (Virtual FAT) extends FAT to support long filenames with Unicode, maintaining backward compatibility.

File sizes are stored as 32-bit integers, meaning the largest file size FAT32 can handle is 4 GB.


File sizes are stored as 32-bit integers.

Volume Organization


In FAT file systems, the storage device is divided into specific regions, each serving a defined role:

- The Volume Boot Record (VBR) contains the so-called BIOS Parameter Block
- The root directory of FAT12/16 has a fixed location and size
- FAT32 boot sector includes the locations of the root directory, FSINFO structure (that keeps track of free clusters, to optimize allocations), and boot-sector backup (should be 6)

FAT12/16



FAT32



1. Reserved Area


- Starts at sector 0.
- Contains the Volume Boot Record (VBR or Boot sector), which holds key information about the file system.
- FAT12/16: Usually 1 sector (only the VBR); FAT32: Larger because it includes FSINFO structure (helps track free clusters).

2. FAT Area

- follows the reserved area, and its size is calculated by multiplying the number of tables by their size

3. Data Area

- Clusters are only in Data Area, numbered from 2 (!!!) and after the root directory for FAT12/16
- Data could be also hidden after the last valid entry in a FAT table



The “Small Sectors” and “Large Sectors” fields represent the total number of sectors in the volume. Only one of these fields is used, and the other is set to zero.

Green (BIOS Parameter Block - BPB):

- Essential fields required for the basic operation of the file system.
- Defines sector sizes, cluster sizes, and disk structure.

Yellow (Extended BIOS Parameter Block - EBPB):

- Additional metadata introduced in later FAT versions.
- Includes details like the Volume Serial Number, Boot Signature, and Volume Label.

Boot sector FAT12/16

FAT16 Boot Sector

	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F		
0	Jump Instruction			OEM ID									Bytes / sect		sect / cluster		reserved sectors	
10	no / FATS	Root entries		Small Sectors		Media descriptor	Sectors / FAT		Sectors / Track		Number / heads		Hidden sectors					
20	large Sectors			Physical drive number	reserved	ext boot sig	Volume Serial Number				Volume Label (deprecated)							
30	Volume label						File System Type											
40	OS Boot Code																	
50																		
60																		
70																		
...																		
1D0																		
1E0																		
1F0															55	AA		

Boot sector FAT32

FAT32 Boot Sector																		
	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F		
0	Jump Instruction			OEM ID									Bytes / sect		sect / cluster		reserved sectors	
10	no / FATS		0x0000		0x0000		Media descriptor	0x0000		Sectors / Track		Number / heads		Hidden sectors				
20	large Sectors Total sectors in volume				Sectors / FAT				0x0000		File System Version		Root(first) Cluster Number					
30	FS Info sector		Backup boot sector		Reserved													
40	Phys Drive num	0x00	Extd boot sig	Volume Serial Number				Volume Label (deprecated) normally "NO NAME "										
50	Vol Label		System ID "FAT32"									Boot code						
60	Boot code																	
70																		
80																		
90																		
...																		
1D0																		
1E0																		
1F0															55	AA		

https://www.writeblocked.org/resources/FAT_cheatsheet.pdf

Files and Directories

A directory entry in FAT file systems is a 32-byte record that stores metadata about a file or directory.

FAT Directory Entry																
	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F
0	File name								Extension			attribute	reserved	10ms create time ¹	create time	
10	create date		last access date		unused		modified time		modified date		start cluster		File Size			

1. The 10millisecond create time is technically only used in FAT32.

The File Allocation Table (FAT) keeps track of file storage using cluster chains. Each file's data is stored in clusters, and the FAT table links these clusters together to form a chain.


- Each FAT entry points to the next cluster in the file's.
- Clusters marked EOF (End of File) indicate the last cluster of a file.

fsstat decodes this chains (in sectors); special values:

0 → not allocated
 0xf...ff0-0xf...f6 → reserved
 0xf...ff7 → damaged
 0xf...ff8-0xf...fff → EOF

Note: FAT entries start at 0, but:

- The first addressable cluster #2.



- Entry 0 typically stores a copy of the media type, and entry 1 stores the dirty-status of the file system

Example

In `eighties.dd` (SHA256: `cc121c3a...`) and `eighties-all-files.dd` (SHA256: `e5f16884...`) you'll find two very similar FAT16 (not VFAT) file systems. In the former all files have been deleted. Using `ImHex`.

1. find out: Sector and cluster sizes Number of reserved sectors Locations of: FAT1, FAT2, Root Dir. (=Data Area), first cluster (#2)

compare these results with the output of `fsstat`

2. check the FAT entries for `48.gif` in the two `dd`-images, and compare the results of `istat` on “inode” 5

```
fls -r eighties-all-files.dd
```

```
##OUT##
d/d 3:  jpgs
+ r/r 517:      clive.jpg
+ r/r 518:      48k.jpg
d/d 4:  games
+ r/r 581:      mmonty.tzx
r/r 5:  48.gif
r/r 6:  48.txt
v/v 523203:     $MBR
v/v 523204:     $FAT1
v/v 523205:     $FAT2
V/V 523206:     $OrphanFiles
###
```

```
istat eighties-all-files.dd 5
```

```
##OUT##
Sectors:
108 109 110 111 116 0 0 0 #l'ha recuperato dentro la FAT
###
```

```
fsstat eighties-all-files.dd
```

```
##OUT##
FAT CONTENTS (in sectors)

100-103 (4) -> EOF
104-107 (4) -> EOF
108-111 (4) -> 116 # cluster chain
112-115 (4) -> EOF
116-119 (4) -> EOF
120-331 (212) -> EOF
332-1211 (880) -> EOF
1212-1279 (68) -> EOF
###
```

When a file name exceeds the 8.3 format, the file system creates additional directory entries to store the name in Unicode (2 bytes per character).


These LFN entries are linked together and precede the main directory entry (which still stores the short 8.3 name for compatibility).

- Each LFN entry is marked with the attribute 0x0F, meaning it is not treated as a normal file entry.
- The last LFN entry in the sequence has its sequence number OR-ed with 0x40; or 0xe5 if unallocated.

Long File Name

	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F	
0	file name (Unicode 2 bytes/char)												0x0F	reserved	Check sum	file name	
10	file name											0x0000		file name			

The quick brown fox", as THEQUI~1.FOX in 8.3 convention.




When a new directory is created, it contains

. and ..

Cluster 110			Cluster 196		
Name	Created	Cluster	Name	Created	Cluster
dir2	3/30/04 01:29:01	128	.	4/1/04 09:27:00	196
dir1	4/03/04 11:47:40	196	..	4/1/04 09:27:00	110
file8.dat	3/30/04 20:41:12	112	file1.dat	4/3/04 12:58:23	297

those entries can be helpful for carving deleted directories



Since the size of a directory is always 0, the only way to know how many cluster to read is following the cluster chain

Example

In eighties-vfat.dd (SHA256: 62258f92ebb42226...) and eighties-vfat-all-files.dd (SHA256: fe46141b98d227cb...) you'll find two very similar, and familiar, VFAT FAT16 file systems. As with the previous exercise, in the former all files have been deleted.

Yet, `fls -rp eighties-vfat.dd` can show the full, long name, for some deleted files but not for others, that are listed under `$OrphanFiles`.

1. Can you explain why? Hint: eighties-vfat-all-files.dd contains some clues

In some cases, even if the file is cancelled, we have the full name, and it is strange because in the fat one byte '_' is put above the first character (eighties.dd). Since we have a vfat there are the entries with the long name and we can trace the original name. OrphanFiles is a standard used by TSK when it does not have babstanz ainomraizons to know where that file is.

2. Using ImHex, can you manually recover the full names from eighties-vfat.dd?

```
fls -rp eighties-vfat.dd
##OUT##
r/r * 3:      -
r/r * 4:      -
d/d * 6:      Games
r/r * 583:    Games/Mutant Monty.tzx
r/r * 7:      _8.gif
r/r * 8:      _8.txt
v/v 523203:   $MBR
v/v 523204:   $FAT1
v/v 523205:   $FAT2
V/V 523206:   $OrphanFiles
-/r * 519:    $OrphanFiles/_LIVES~1.JPG
-/r * 520:    $OrphanFiles/_8k.jpg
###
```

NTFS

New Technology File System (NTFS) is a proprietary journaling file system developed by Microsoft in 1993, supporting:

- Access Control Lists (ACLs)
- Encryption
- Transparent compression
- Sparse files
- Journaling
- POSIX support (no, not WSL)
- Multiple data streams

A nice feature is that everything is a file!

- Except for the Volume Boot Record (VBR), everything else is considered a data area.
- Any sector (except VBR) can be allocated to a file.
- A very scalable design where internal structures can change over time.
- Generic data structures embed specific content.

Each unit of information associated with a file is implemented as a file attribute (NTFS Object Attribute). Each attribute consists of a byte stream.

Key points about NTFS attributes:

- The contents of a file are treated as 'an attribute,' similar to its name or timestamps.
- Each file has a special \$DATA attribute with no name that corresponds to its content.

Applications can create additional named streams, called Alternate Data Streams (ADS).

- For example, the \$Zone.Identifier is used by Windows to mark files downloaded from the web.

- You can list ADSs using `dir /r` or tools like `streams`.
- You can set or show their contents by redirecting commands like `echo` and `more`.

When mounting NTFS file systems, consider specifying:

- `show_sys_files` to display metafiles in directory listings.
- `streams_interface=windows` to access ADS like in Windows.

Refer to `mount.ntfs(8)` for more details.

Hard links allow multiple paths to refer to the same file (not directory):

```
mklink /h new-name existing-name
```

- Similar to Unix, they are reference-counted and limited to the same file system.

Soft links are strings interpreted dynamically and can point to files, directories, or non-existent targets:

```
mklink new-name existing-name
```

- Implemented as reparse points (files or directories containing application-specific reparse data and a 32-bit reparse tag).

Junctions are a legacy concept, functioning almost identically to directory symbolic links. See [ARIS21] for more details.

Shortcuts are `.lnk` files interpreted by Explorer.

Volume Organization

NTFS maps the entire volume into **clusters**, which are the smallest allocatable units of storage.

This is different from FAT, which uses a simpler allocation table.

The **cluster factor** (size of a cluster) depends on the size of the volume.

It is an integral number of physical sectors and always a power of 2.

- Larger clusters:
 - Reduce fragmentation.
 - Speed up allocation.
 - However, they can lead to *wasted space* (internal fragmentation) when storing small files.

1. Logical Cluster Numbers (LCNs)

LCNs are used to number clusters sequentially from the beginning of the volume.

They represent the physical location of clusters on the disk.

2. Virtual Cluster Numbers (VCNs)

VCNs are used to address data within a file.

They represent the logical view of a file's data.

VCNs are not necessarily physically contiguous, allowing NTFS to handle fragmented files efficiently.

Volume Layout

0–2	Assembly instruction to jump to boot code	No (unless it is the bootable file system)
3–10	OEM Name	No
11–12	Bytes per sector	Yes
13–13	Sectors per cluster	Yes
14–15	Reserved sectors (Microsoft says it must be 0)	No
16–20	Unused (Microsoft says it must be 0)	No
21–21	Media descriptor	No
22–23	Unused (Microsoft says it must be 0)	No
24–31	Unused (Microsoft says it is not checked)	No
32–35	Unused (Microsoft says it must be 0)	No
36–39	Unused (Microsoft says it is not checked)	No
40–47	Total sectors in file system	Yes
48–55	Starting cluster address of MFT	Yes
56–63	Starting cluster address of MFT Mirror \$DATA attribute	No
64–64	Size of file record (MFT entry)	Yes
65–67	Unused	No
68–68	Size of index record	Yes
69–71	Unused	No
72–79	Serial number	No
80–83	Unused	No
84–509	Boot code	No

8-

1. No Fixed Layout:

- NTFS does not have a fixed layout for the volume, except for the Volume Boot Record (VBR).
- The VBR contains information that guides the system to the [s](#), which is the core of NTFS.

2. 8-bit Sizes Interpretation:

- Sizes in NTFS are represented as 8-bit values.
- Positive values indicate the number of clusters.
- Negative values are interpreted as 2^n , where n is the negative value.

Master File Table The MFT is the heart of the NTFS volume structure. It is implemented as an array of file records (similar to inodes). Key details include:

- The size of each record can be defined at format time (typically 1 KB or 4 KB).
- The size depends on the underlying physical medium:
 - Disks with 4 KB sector sizes generally use 4 KB file records.
 - Older disks with 512-byte sectors use 1 KB file records.
- The size does not depend on the cluster size.

When a file needs more metadata space, the base file record stores the location of additional records. The MFT's location

is specified in the BIOS Parameter Block inside the VBR.

Key points about the MFT:

- The MFT contains one record for each file, including itself (the first entry).
- It can be fragmented; however, an MFT zone is typically reserved during formatting (about 12.5% of the entire volume).
- In addition to file records, there are other system metadata files. These hidden files have names that begin with a dollar sign (\$).

File records for metadata files

0	\$MFT - MFT	Reserved for NTFS metadata files
1	\$MFTMirr - MFT mirror	
2	\$LogFile - Log file	
3	\$Volume - Volume file	
4	\$AttrDef - Attribute definition table	
5	\ - Root directory	
6	\$Bitmap - Volume cluster allocation file	
7	\$Boot - Boot sector	
8	\$BadClus - Bad-cluster file	
9	\$Secure - Security settings file	
10	\$UpCase - Uppercase character mapping	
11	\$Extend - Extended metadata directory	
12	Unused	
23	Unused	
24	\$Extend\ \$Quota - Quota information	
25	\$Extend\ \$ObjId - Distributed link tracking information	
26	\$Extend\ \$Reparse - Back references to reparse points	
27	\$Extend\ \$RmMetadata - RM metadata directory	
28	\$Extend\ \$RmMetadata\ \$Repair - RM repair information	
29	\$Extend\ \$Deleted - POSIX deleted files	
30	\$Extend\ \$RmMetadata\ \$TxflLog - TxF log directory	
31	\$Extend\ \$RmMetadata\ \$Txfl - TxF metadata directory	
32	\$Extend\ \$RmMetadata\ \$TxflLog\ \$Tops - TOPS file	
33	\$Extend\ \$RmMetadata\ \$TxflLog\ \$TxflLog.blf - TxF BLF	
34	\$TxflLogContainer000000000000000000000001	
35	\$TxflLogContainer000000000000000000000002	

NTFS incorporates [fixup values](#) into data structures that are over one sector in length

- the last two bytes of each sector are replaced with a “signature” value

- the signature is later used to verify the integrity of the data

Fixups are only in data structures, not in sectors with file content

The “signature” is incremented each time the structure is updated

File identifiers

Files identified by 64-bit “file record numbers”, which consist of:

1. a file number, corresponding to the (0-based) position in the MFT
2. a sequence number, incremented when a file record is reused



MFT Entries

File records start with a fixed header:

File Record Segment Header																
	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F
0	F	I	L	E	Update Seq array offset		Update Seq array size		\$LogFile Sequence Number							
1	Seq no		Hard Link Count		1 st attrib offset		Flags		Used size of file record				Allocated size of file record			
2	File reference to base file record								Next attrib ID				MFT Record No			

followed by attributes (and fixup values):

(‘Update Seq == fix-up) NTFS Cheat Sheets



EXT

The Ext (Extended) file system family—Ext2, Ext3, and Ext4—is a series of filesystems used in Linux. These are based on the traditional Unix File System (UFS), but add several modern capabilities over time.

Key Modern Features:

- Extended Attributes / POSIX ACLs: Allow more detailed permissions than traditional Unix read/write/execute.
- Journaling (from Ext3 onwards): Keeps a log (journal) of changes to help recover quickly after crashes.
- Encryption: Protects file contents and filenames (introduced in Ext4).
- 64-bit Support: Allows management of much larger files and filesystems.
- More...

The Ext file systems are modular, with features grouped by their compatibility level:

Feature Type	What It Means	Example
Compatible	If not supported, OS can still mount the FS normally	has_journal
Incompatible	If not supported, OS must not mount the FS	encrypt
Read-Only Compatible	If not supported, OS can mount in read-only mode	dir_index

Support for Very Large Volumes

- Ext3: Limited to volumes up to 16 TB (terabytes), calculated as $16 \cdot 10^{12} / 2^{44}$.
- Ext4: Extends this limit to 1 exabyte (EB), or $10^{18} / 2^{60}$. This makes Ext4 suitable for modern systems with very high storage demands.

Maximum File Size

- Ext3: The maximum size of a single file is 2 TB.
- Ext4: Increases this limit to 16 TB, making it more suitable for applications requiring the management of large files.

Backward Compatibility

- Ext4: Can mount Ext3 file systems and treat them as Ext4. This process is called “on-the-fly conversion”: existing files remain in the Ext3 format, but new files will use Ext4’s advanced features.
- Ext3: Cannot mount Ext4 file systems, as Ext3 does not support the new features introduced in Ext4.

Extents

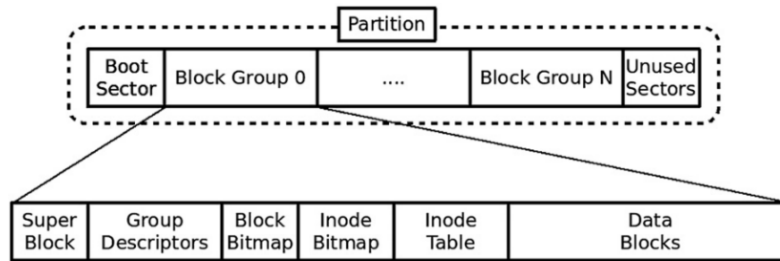
Ext4 uses a structure called extents to map data blocks more efficiently. Extents represent a continuous range of blocks, reducing fragmentation and improving performance. This approach is similar to the one used by NTFS (the Windows file system).

Layout

In the Ext file system family (Ext2, Ext3, Ext4):

- Sectors (the smallest physical units on a disk, typically 512 bytes) are grouped into **blocks** of 1, 2, or 4 KB.
- These blocks are conceptually similar to “clusters” in FAT or NTFS.
- Unlike the Unix File System (UFS), Ext does not assemble blocks from smaller “fragments.” In Ext, the terms “block” and “fragment” are essentially synonymous (the term “block” is predominantly used).
- The **superblock** is located at an offset of 1024 bytes from the start of the volume (not at the very beginning). It contains all the fundamental parameters of the file system.
- The size of the superblock is 1024 bytes. An optional reserved area may follow the superblock.

- Boot code, if present, resides in the Master Boot Record (MBR) or Volume Boot Record (VBR), not within the file system itself.
- The rest of the file system is divided into **block groups**.



A key detail: when the first block group (BG0) starts at offset 0:


- The primary superblock is located at an offset of 1024 bytes from the volume start (or within block 0).
- Backup copies of the superblock (if present) are located at the beginning of other block groups (BG1, BG2, etc.).

Each “copy” of the superblock has its own specific metadata.

- This backup structure, with superblock copies in multiple locations, is a critical safety feature that allows recovery even if parts of the file system are damaged.

Each block group

1. contains its data bitmap, inode bitmap, inode table, data blocks
2. may contain the backup of the superblock and the GD table



Each **group descriptor** defines where to find the components of the block group.

- The main **group descriptor table** is located in the block immediately following the superblock.
- This table can be copied, after the superblock backup, to all or some groups, depending on the formatting options (see `sparse_super[2]` in the `ext4(5)` man page).

Ext4 introduces advanced options for managing metadata and resizing:

- **flexible block groups (flex_bg)**: allows block group metadata to be placed anywhere, not necessarily within the same group.
- **reserved GDT blocks**: enables online resizing of the file system without requiring it to be unmounted.

Standard command `dumpe2fs` prints superblock and group information.

- All groups, except for the last, contain the same number of blocks n
- by default, n is equal to the # of bits in a block: per blocchi di 4 KB (4096 byte), ogni blocco contiene 32.768 bit (4096×8), quindi ogni gruppo può gestire 32.768 blocchi
 - the block bitmap is exactly 1 block
- the # of inodes can be chosen at format time, usually less than n

Example

Use `dumpe2fs/fsstat` to list

- block size = 2048 byte
- blocks and inodes per group = 888;296
- number of groups = 18

in the image `linux-ext2.dd` (SHA256: `b6b1836ff1efef3a70...`)

- Are there unused sector after the FS? no total size of fs == $16384 * 2048$ == `stat -c %s linux-ext2.dd`
- All are group of the same size? no, the end group have a size of 400 blocks

Unused space

- The first 1024 bytes are not technically used
- There are unused bytes in the superblock
- There can be unused entries in the GD-table
 - Or they backup copies
- There are the reserved GDT blocks

Those are possible places where to hide data

Superblock details

Byte Range	Description	Essential
0–3	Number of inodes in file system	Yes
4–7	Number of blocks in file system	Yes
8–11	Number of blocks reserved to prevent file system from filling up	No
12–15	Number of unallocated blocks	No
16–19	Number of unallocated inodes	No
20–23	Block where block group 0 starts	Yes
24–27	Block size (saved as the number of places to shift 1,024 to the left)	Yes
28–31	Fragment size (saved as the number of bits to shift 1,024 to the left)	Yes
32–35	Number of blocks in each block group	Yes
36–39	Number of fragments in each block group	Yes
40–43	Number of inodes in each block group	Yes
44–47	Last mount time	No
48–51	Last written time	No
52–53	Current mount count	No
54–55	Maximum mount count	No
56–57	Signature (0xef53)	No
58–59	File system state (see Table 15.2)	No
60–61	Error handling method (see Table 15.3)	No
62–63	Minor version	No
64–67	Last consistency check time	No
68–71	Interval between forced consistency checks	No
72–75	Creator OS (see Table 15.4)	No
76–79	Major version (see Table 15.5)	Yes
80–81	UID that can use reserved blocks	No
82–83	GID that can use reserved blocks	No
84–87	First non-reserved inode in file system	No
88–89	Size of each inode structure	Yes

Byte Range	Description	Essential
90–91	Block group that this superblock is part of (if backup copy)	No
92–95	Compatible feature flags (see Table 15.6)	No
96–99	Incompatible feature flags (see Table 15.7)	Yes
100–103	Read only feature flags (see Table 15.8)	No
104–119	File system ID	No
120–135	Volume name	No
136–199	Path where last mounted on	No
200–203	Algorithm usage bitmap	No
204–204	Number of blocks to preallocate for files	No
205–205	Number of blocks to preallocate for directories	No
206–207	Unused	No
208–223	Journal ID	No
224–227	Journal inode	No
228–231	Journal device	No
232–235	Head of orphan inode list	No
236–1023	Unused	No

Example

1. use `dumpe2fs/fsstat` to check where the superblock copies are
2. dump the first three, and compare them: are they equal?
 - note: the main SB starts at offset 1024, the others at 0
 - are the 2nd and 3rd equal?

```
dumpe2fs linux-ext2.dd | grep "superblock"
```

```
##OUT##
```

```
Primary superblock at 0, Group descriptors at 1-1
```

```
Backup superblock at 888, Group descriptors at 889-889
```

```
Backup superblock at 2664, Group descriptors at 2665-2665
```

```
Backup superblock at 4440, Group descriptors at 4441-4441
```

```
Backup superblock at 6216, Group descriptors at 6217-6217
```

```
Backup superblock at 7992, Group descriptors at 7993-7993
```

```
###
```

```
#The primary superblock is located after the boot block
```

```
# the superblock, at volume offset 1024, contains the FS parameters
```

```
# the superblock is 1024 bytes in size
```

```
dd if=linux-ext2.dd bs=1k skip=1 count=1 of=sb1
```

```
#To convert from filesystem block (2048) numbers to 1KB blocks for dd, you need to multiply by 2:
```

```
dd if=linux-ext2.dd bs=1k skip=$((888*2)) count=1 of=sb2
```

```
dd if=linux-ext2.dd bs=1k skip=$((2664*2)) count=1 of=sb3
```

```
md5sum sb1 sb2 sb3
```

```
##OUT##
```

```
922c7935bd45a7b36fc938f3b000d16c sb1
```

```
06469d508b7da81f55e090addf908598 sb2
```

```
5858141cac833da5d7a0195dfd968119 sb3
```

```
###
```

```
vbindiff sb2 sb3
```

```
###
change the bit that keeps track of the superblock group membership
###
```

```
vbindiff sb1 sb3
```

```
###
more diversity because the primary superblock also times the last path where the fs was mounted (no
###
```

Group descriptors details

Byte Range	Description	Essential
0–3	Starting block address of block bitmap	Yes
4–7	Starting block address of inode bitmap	Yes
8–11	Starting block address of inode table	Yes
12–13	Number of unallocated blocks in group	No
14–15	Number of unallocated inodes in group	No
16–17	Number of directories in group	No
18–31	Unused	No

Inode

Inodes are the primary metadata structure.

- fixed size, defined in the superblock, minimum 128 bytes (`dumpe2fs -h linux-ext2.dd | grep "Inode size"`)
- extra-space can be used to store extended attributes (otherwise, data blocks are used)
- small non-user content can be stored inside direct-block pointer array; e.g., symlink values
- each inode has a unique ‘address’ or number, starting with 1 (index 0 is not used)
- The inodes of each block group are stored in a table and the position of this table is specified in the block descriptor
- The allocation status of each inode is determined via the inode bitmap and the position of this table is specified in the block descriptor

Standard commands can be used to inspect the FS:

- `-i,-lai` in `ls`, shows inode-numbers
- `stat` shows some metadata

`debugfs` is a FS “debugger” that can display a lot of information.

```
ln pippo.txt pippo-hardlink.txt
ln -s pippo.txt pippo-hardlink.txt
ls -li
##OUT##
1445743 -rw-r--r--  2 root   root      0 Apr  3 10:20 pippo-hardlink.txt
1445746 lrwxrwxrwx  1 root   root      9 Apr  3 10:42 pippo-symlink.txt -> pippo.txt
1445743 -rw-r--r--  2 root   root      0 Apr  3 10:20 pippo.txt
###
```

###

for symb link we get different inode number because is the newfile and fs must allocate a new inode and content for symb link. The content of ippon-symblink.txt is the path " -> pippo.txt", that is just a string that is resolved when we try to open the ippon-symblink.txt.

###

Inodes 1 to 10 are reserved

- 1 is (was?) used for keeping track of bad blocks
- 2 is used for the root directory
- 8 is typically for the journal


Byte Range	Description	Essential
0-1	File mode (type and permissions)	Yes
2-3	Lower 16 bits of user ID	No
4-7	Lower 32 bits of size in bytes	Yes
8-11	Access Time	No
12-15	Change Time	No
16-19	Modification time	No
20-23	Deletion time	No
24-25	Lower 16 bits of group ID	No
26-27	Link count	No
28-31	Sector count	No
32-35	Flags	No
36-39	Unused	No
40-87	12 direct block pointers	Yes
88-91	1 single indirect block pointer	Yes
92-95	1 double indirect block pointer	Yes
96-99	1 triple indirect block pointer	Yes
100-103	Generation number (NFS)	No
104-107	Extended attribute block (File ACL)	No
108-111	Upper 32 bits of size / Directory ACL	No
112-115	Block address of fragment	No
116-116	Fragment index in block	No
117-117	Fragment size	No
118-119	Unused	No
120-121	Upper 16 bits of user ID	No
122-123	Upper 16 bits of group ID	No
124-127	Unused	No

Directories


The old directory entry was a simple data structure containing

- the file name (variable length, 1-255 chars) where the length of an entry is rounded up to a multiple of four
- the inode number (AKA index+1)

A)



B)



The newer one uses one byte in the filename-length to store the file type.

Allows detecting when an inode has been reassigned (reallocated)!

Byte Range	Description	Essential
0–3	Inode value	Yes
4–5	Length of this entry	Yes
6–6	Name length	Yes
7–7	File type	No
8+	Name in ASCII	Yes

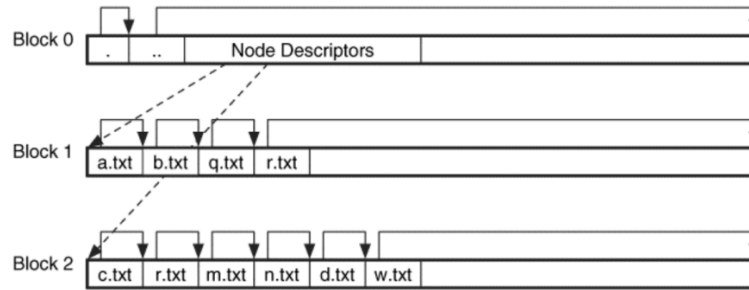
File Type	Description
1	Regular file
2	Directory
3	Character device
4	Block device
5	FIFO
6	Unix Socket
7	Symbolic link

In ext3 and ext4, option `dir_index` allows the FS to use hashed B-trees to speed up name lookups

- similar to NTFS: This approach is conceptually similar to that used by NTFS in Windows
- when this functionality is active, a `ro-compatible` (read-only compatible) flag is set in the superblock (This means that systems that do not support this function can still mount the file system, but only in read-only mode)

Still the same directory-entry structures, but in sorted order

- each entry continues to have the same format and contain the same information (Inode value, Length of this entry ...) but instead of storing the entries in a linear manner (one after the other), the entries are organised in a B-tree structure



Example

List files in the root directory of linus-ext2.dd by

- (ro) mounting it
- using TSK's fls
- using ImHex
 - get the block-size from the SB == 1 (=1K)
 - the GD-table is in the block after the SB
 - decode the first GD and get the block for the (first) inode-table
 - the inode n.2 (index 1) is for the root directory
 - the first direct-block-pointer points to the beginning of the directory
 - decode the directory entries

```
mount -oro linus-ext2.dd mnt/linus-ext2
```

```
ls -li
```

```
##OUT##
```

```
total 28
```

```
12 -rw-r----- 2 root root 773 Apr 8 2023 hard_link.txt
```

```
12 -rw-r----- 2 root root 773 Apr 8 2023 linus.txt
```

```
15 lrwxrwxrwx 1 root root 109 Apr 9 2023 long_symlink -> /media/someuser/whatever/subdir1/subdi
```

```
11 drwx----- 2 root root 16384 Apr 8 2023 lost+found
```

```
5033 drwxr-xr-x 4 root root 2048 Apr 8 2023 pics
```

```
13 lrwxrwxrwx 1 root root 4 Apr 8 2023 pictures -> pics
```

```
14 -rw-r--r-- 1 root root 8 Apr 9 2023 very_short_text.txt
```

```
###
```

```
fls -rp linus-ext2.dd
```

```
##OUT##
```

```
d/d 11: lost+found
```

```
d/d 5033: pics
```

```
d/d 5034: pics/tux
```

```
r/r 5040: pics/tux/tux-lego.png
```

```
r/r 5041: pics/tux/tux-windows.png
```

```
r/r 5042: pics/tux/tux.png
```

```
d/d 5035: pics/linus
```

```
r/r 5036: pics/linus/linus_2018.jpg
```

```
r/r 12: linus.txt
```

```
l/l 13: pictures
```

```
r/r 12: hard_link.txt
```

```
r/r 14: very_short_text.txt
```

```

l/l 15: long_symlink
V/V 5625:      $OrphanFiles
-/r * 5037:    $OrphanFiles/OrphanFile-5037
-/r * 5038:    $OrphanFiles/OrphanFile-5038
-/r * 5039:    $OrphanFiles/OrphanFile-5039
###

```

Extended

Extended attributes are an advanced feature of Ext file systems (Ext2/3/4) that allows additional metadata to be associated with files and directories. This metadata is stored as name-value pairs.

- Ogni attributo è una coppia composta da:
 - Name: specified in the `namespace.attribute-name` format (current namespaces are security, system, trusted, and user)
 - Value: An arbitrary datum associated with the name

Gli attributi estesi sono utilizzati per implementare le Access Control Lists (ACL), che consentono di definire permessi più dettagliati rispetto ai tradizionali permessi Unix.

```

setfacl -m u:username:rwX file.txt # Set up an ACL
getfacl file.txt                   # View ACLs

```

kernel and filesystem may place limits on the maximum number and size of extended attributes that can be associated with a file

Extended attributes are stored in a separate block associated with the file. If several files share the same attributes, they can use the same block to optimise space.

```

getfattr -d file.txt # Display extended attributes

```

```

# set an extended attribute
setfattr -n user.comment -v "This is an important file" file.txt
setfattr -x user.comment file.txt # setfattr -x user.comment file.txt

```

See `xattr(7)`