

This work is licensed under a Creative Commons license



Attribution-NonCommercial-NoDerivatives 4.0 International  
(CC BY-NC-ND 4.0)

You are free to:

**Share** copy and redistribute the material in any medium or format.

Under the following terms:

**Attribution** You must give appropriate credit, provide a link to the license, and indicate if changes were made. You may do so in any reasonable manner, but not in any way that suggests the licensor endorses you or your use.

**NonCommercial** You may not use the material for commercial purposes.

**NoDerivatives** If you remix, transform, or build upon the material, you may not distribute the modified material.

# The FAT File System Family

Giovanni Lagorio

`giovanni.lagorio@unige.it`  
`https://csec.it/people/giovanni\_lagorio`  
`X/GitHub/...: zxgio`


DIBRIS - Dipartimento di Informatica, Bioingegneria, Robotica e Ingegneria dei Sistemi  
University of Genova, Italy

# Outline

- 1 The FAT family
- 2 Volume Organization
- 3 Files and Directories

# File System Forensics Analysis

The “bible” for this part of the course is Brian Carrier’s “**File System Forensics Analysis**” [Car05], where you can find more details



# Introduction

First versions of FAT were developed in 1977/1978 (!), FAT12 in 1980

- No ACLs, no quota, no journal
- Still very common; lightweight and suited for embedded solutions

Three versions: **FAT12**, **FAT16** and **FAT32**

- The number indicates the **# of bits used to identify clusters**
  - FAT12 can address  $2^{12} = 4096$  clusters. Windows permits cluster sizes from 512 bytes to 8 KB, which limits FAT12 to 32 MB
  - FAT16 can address  $2^{16} = 65,536$  clusters
  - FAT32 can address  $2^{28}$  clusters (top 4 bits used for other purposes)

Actually, **first 2 & last 16 are reserved**: usable clusters are slightly less

- **Timestamps are in local time**
- **Filenames are in MSDOS 8.3 format**
  - VFAT: a backward compatible extension allowing Unicode long names
- File **sizes are stored as 32-bit integers**

## exFAT/FAT64

- introduced in 2006, but specification published in 2019
- some **advanced functionality** (e.g. optional ACLs), can be used where NTFS is unfeasible
- a bitmap tracks free clusters, **improving the performance** of allocation and deletion operations
- specifically **designed for flash drives**
- as the name implies, file **sizes stored as 64 bit integers**

We'll study the “classic” FAT family


# Outline

- 1 The FAT family
- 2 Volume Organization
- 3 Files and Directories


# Organization

- The VBR contains the so-called **BIOS Parameter Block**
- The **root directory of FAT12/16** has a fixed location and size
- **FAT32** boot sector includes the locations of the **root directory**, **FSINFO** structure (that keeps track of free clusters, to optimize allocations), and **boot-sector backup** (should be 6)

FAT12/16



FAT32





# Boot sector FAT12/16

**FAT16 Boot Sector**

FAT16 Boot Sector																		
	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F		
0	Jump Instruction			OEM ID									Bytes / sect		sect / cluster		reserved sectors	
10	no / FATS	Root entries		Small Sectors		Media descriptor	Sectors / FAT		Sectors / Track		Number / heads		Hidden sectors					
20	large Sectors			Physical drive number	reserved	ext boot sig	Volume Serial Number			Volume Label (deprecated)								
30	Volume label						File System Type											
40	OS Boot Code																	
50																		
60																		
70																		
...																		
1D0																		
1E0																		
1F0																		

[https://www.writeblocked.org/resources/FAT\\_cheatsheet.pdf](https://www.writeblocked.org/resources/FAT_cheatsheet.pdf)

“small sectors”/“large sectors”=“# of sectors”: only one is used, the other is 0  
In green the **Bios Parameter Block**, in yellow the **BPB extended**

*Root entries* and *Sectors/FAT* (AKA size in sectors) are 0 for FAT32...

# Boot sector FAT32

**FAT32 Boot Sector**

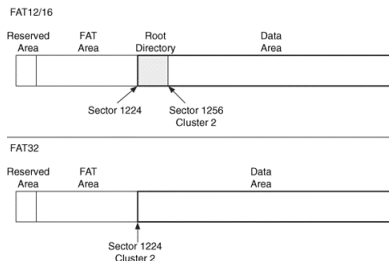
FAT32 Boot Sector																		
	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F		
0	Jump Instruction			OEM ID									Bytes / sect		sect / cluster		reserved sectors	
10	no / FATS	0x0000		0x0000		Media descriptor	0x0000		Sectors / Track		Number / heads		Hidden sectors					
20	large Sectors Total sectors in volume				Sectors / FAT				0x0000		File System Version		Root(first) Cluster Number					
30	FS Info sector		Backup boot sector		Reserved													
40	Phys Drive num	0x00	Extd boot sig	Volume Serial Number				Volume Label (deprecated) normally "NO NAME "										
50	Vol Label		System ID "FAT32"									Boot code						
60	Boot code																	
70																		
80																		
90																		
...																		
1D0																		
1E0																		
1F0															55	AA		

[https://www.writeblocked.org/resources/FAT\\_cheatsheet.pdf](https://www.writeblocked.org/resources/FAT_cheatsheet.pdf)

# How to locate the areas


- 1 The **reserved area** starts at sector 0, its size is in the VBR (1 means only the VBR is reserved)
  - Usually 1 for FAT12/16; FAT32 uses more because of FSINFO
- 2 **FAT area** follows the reserved area, and its size is calculated by multiplying the number of tables by their size
- 3 then, there is the **Data Area**

**Clusters are only in Data Area, numbered from 2 (!!!) and after the root directory for FAT12/16**



# Final sectors

The size of the data area may not be a multiple of the cluster size, so there could be unused sectors at the end:



Data could be also hidden after the last valid entry in a FAT table

# Outline

- 1 The FAT family
- 2 Volume Organization
- 3 Files and Directories

# Directory entry

A directory entry is 32 bytes and stores

- **file name (8.3)** — first byte 0xe5 means deleted (0x05 used to encode 0xe5)
- **attributes**: RO 1, hidden 2, sys 4, vol. label 8, dir 16, archived 32
- **size, a 32-bit integer** (0 for directories)
- **starting cluster**
- **time stamp information**
  - created, last accessed, and last written with widely different granularity
  - created/access timestamps are optional


**FAT Directory Entry**

	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F
0	File name								Extension			attribute	reserved	10ms create time <sup>1</sup>	create time	
10	create date		last access date		unused		modified time		modified date		start cluster		File Size			

1. The 10millisecond create time is technically only used in FAT32.

[https://www.writeblocked.org/resources/FAT\\_cheatsheet.pdf](https://www.writeblocked.org/resources/FAT_cheatsheet.pdf)

# Cluster chains



`fsstat` decodes this chains (in *sectors*); special values:

- 0 → not allocated
- 0xf...ff0-0xf...f6 → reserved
- 0xf...ff7 → damaged
- 0xf...ff8-0xf...fff → EOF

**FAT entries start at 0, but...**

The first addressable cluster #2. Entry 0 typically stores a copy of the media type, and entry 1 stores the dirty-status of the file system

## Demo/Exercise

In `eighties.dd` (SHA256: `cc121c3a...`) and `eighties-all-files.dd` (SHA256: `e5f16884...`) you'll find two very similar FAT16 (not VFAT) file systems. In the former all files have been deleted.

Using ImHex,

① find out:

- Sector and cluster sizes
- Number of reserved sectors
- Locations of: FAT1, FAT2, Root Dir. (=Data Area), first cluster (#2)

compare these results with the output of `fsstat`

② check the FAT entries for `48.gif` in the two dd-images, and compare the results of `istat` on “inode” 5



# VFAT/Long File Name support

When a file name does not follow MSDOS 8.3 ASCII convention

- additional entries (with “attributes” 0xf) can store the long name
- those precede the main entry, which stores all other metadata

Long File Name


Long File Name																	
	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F	
0	file name (Unicode 2 bytes/char)												0x0F	reserved	Check sum	file name	
10	file name										0x0000		file name				

[https://www.writeblocked.org/resources/FAT\\_cheatsheet.pdf](https://www.writeblocked.org/resources/FAT_cheatsheet.pdf)

Byte 0 is actually a **sequence number** (the final one is ORed with 0x40) or 0xe5 if unallocated; for example...

# LFN example

“The quick brown.fox”, as THEQUI~1.FOX in 8.3 convention.




From [ARIS21]


# Standard directory entries

When a new directory is created, it contains

- . and ..



- those entries can be helpful for carving deleted directories



Since the size of a directory is always 0, the only way to know how many cluster to read is following the cluster chain

## Demo/Exercise

In `eighties-vfat.dd` (SHA256: 62258f92ebb42226...) and `eighties-vfat-all-files.dd` (SHA256: fe46141b98d227cb...) you'll find two very similar, and familiar, VFAT FAT16 file systems.

As with the previous exercise, in the former all files have been deleted. Yet, `fls -rp eighty-vfat.dd` can show the full, long name, for *some deleted files* but not for others, that are listed under `$OrphanFiles`.

- 1 Can you explain why?

Hint: `eighties-vfat-all-files.dd` contains some clues

- 2 Using `ImHex`, can you manually recover the full names from `eighties-vfat.dd`?

- [ARIS21] Andrea Allievi, Mark Russinovich, Alex Ionescu, and David Solomon.  
*Windows Internals, Part 2, 7th Edition.*  
Microsoft Press, 2021.
- [Car05] Brian Carrier.  
*File System Forensic Analysis.*  
Addison-Wesley Professional, 2005.