

This work is licensed under a Creative Commons license



Attribution-NonCommercial-NoDerivatives 4.0 International
(CC BY-NC-ND 4.0)

You are free to:

Share copy and redistribute the material in any medium or format.

Under the following terms:

Attribution You must give appropriate credit, provide a link to the license, and indicate if changes were made. You may do so in any reasonable manner, but not in any way that suggests the licensor endorses you or your use.

NonCommercial You may not use the material for commercial purposes.

NoDerivatives If you remix, transform, or build upon the material, you may not distribute the modified material.

The Ext File System Family

ext2/ext3 and ext4

Giovanni Lagorio

`giovanni.lagorio@unige.it`
`https://csec.it/people/giovanni_lagorio`
`X/GitHub/...: zxgio`


DIBRIS - Dipartimento di Informatica, Bioingegneria, Robotica e Ingegneria dei Sistemi
University of Genova, Italy

Outline

- 1 Introduction
- 2 Layout
- 3 Inodes
- 4 Directories
- 5 Extended attributes

File System Forensics Analysis

The “bible” for this part of the course is Brian Carrier’s “**File System Forensics Analysis**” [Car05], where you can find more details




However, for the latest additions, see **ext4(5)** and [Fai12]

Introduction

Linux Ext2/3/4

- based on the traditional Unix FS (UFS)



- with modern features
 - extended-attributes/POSIX ACLs
 - journaling
 - encryption
 - 64-bit support
 - ...

Ext-family designed to be **extensible** via optional features, organized into:

compatible if not supported, the OS can still mount the FS

- E.g., **has_journal**

incompatible if not supported, the FS should not be mounted

- E.g., **encrypt** (only filenames and data blocks; dm-crypt is usually better for single-user PCs)

r.o. compatible if not supported, the OS can still r.o. mount the FS

- E.g., **dir_index**

- **support for very large volumes:** ext3 limited to $16\text{TB} \approx 16 \cdot 10^{12}/2^{44}$, ext4 to 1 exabyte $\approx 10^{18}/2^{60}$
 - max file size from 2 TB to 16 TB
- **backward compatibility:** ext4 can mount ext3 file-systems, as ext4 (“converted on-the-fly”: new files will use new features)
 - but ext3 cannot mount ext4
- **extents for more efficient data-block mapping**, like NTFS
- ...


see [Fai12]

Outline

- 1 Introduction
- 2 Layout**
- 3 Inodes
- 4 Directories
- 5 Extended attributes

Blocks

- sectors are grouped into 1/2/4 KB **blocks** (=FAT/NTFS clusters)
 - ext N , differently from UFS, do not assemble blocks starting from “fragments”, so *block* and *fragment* can be seen as synonyms (we’ll use only “block” from now on)
- the **superblock**, at volume offset 1024, contains the FS parameters
 - the superblock is 1024 bytes in size
 - then there is an optional reserved area
- the boot-code, if any, is contained in the MBR/VBR
- the remainder of the FS is divided into **block groups**




from [Fai12]

Beware: when BG0 starts at 0, SB @ offset 1024 of the volume/block 0, its copies (where present) are at the beginning of BGs 1, 2, ... note that each “copy” has its own metadata

Block groups

Each block group

- 1 contains its data bitmap, inode bitmap, inode table, data blocks
- 2 may contain the backup of the superblock and the GD table



- each **group descriptor** defines where to find (1)
- the main **group descriptor table** is located **in the block after the SB**
 - can be copied, after the superblock backup, in all/some groups (depending on format options; see `sparse_super[2]` in `ext4(5)`)
 - `ext4` offers `flex_bg` option, which allows block-group metadata to be placed anywhere
- there also are **reserved GDT blocks** to allow *on-line* resizing of the FS

Non-TSK commands

Standard command `dumpe2fs` prints superblock and group information

- All groups, except for the last, contain the same number of blocks n
- by default, n is equal to the # of bits in a block
 - the `block bitmap` is exactly 1 block
- the # of inodes can be chosen at format time, usually less than n

Demo/exercise

Use `dumpe2fs/fsstat` to list

- block size
- blocks and inodes per group
- number of groups

in the image `linus-ext2.dd` (SHA256: b6b1836ff1efef3a70...)

Are there unused sector after the FS? All are group of the same size?

Unused space

- The first 1024 bytes are not technically used
- There are unused bytes in the superblock
- There can be unused entries in the GD-table
 - Or they backup copies
- There are the *reserved GDT blocks*

Those are possible places where to hide data

Superblock: details (1/3)

Table 15.1. Data structure for the ExtX superblock.

| Byte Range | Description | Essential |
|------------|--|-----------|
| 0–3 | Number of inodes in file system | Yes |
| 4–7 | Number of blocks in file system | Yes |
| 8–11 | Number of blocks reserved to prevent file system from filling up | No |
| 12–15 | Number of unallocated blocks | No |
| 16–19 | Number of unallocated inodes | No |
| 20–23 | Block where block group 0 starts | Yes |
| 24–27 | Block size (saved as the number of places to shift 1,024 to the left) | Yes |
| 28–31 | Fragment size (saved as the number of bits to shift 1,024 to the left) | Yes |
| 32–35 | Number of blocks in each block group | Yes |
| 36–39 | Number of fragments in each block group | Yes |
| 40–43 | Number of inodes in each block group | Yes |
| 44–47 | Last mount time | No |
| 48–51 | Last written time | No |

Superblock: details (2/3)

| | | |
|-------|--|-----|
| 52–53 | Current mount count | No |
| 54–55 | Maximum mount count | No |
| 56–57 | Signature (0xef53) | No |
| 58–59 | File system state (see Table 15.2) | No |
| 60–61 | Error handling method (see Table 15.3) | No |
| 62–63 | Minor version | No |
| 64–67 | Last consistency check time | No |
| 68–71 | Interval between forced consistency checks | No |
| 72–75 | Creator OS (see Table 15.4) | No |
| 76–79 | Major version (see Table 15.5) | Yes |
| 80–81 | UID that can use reserved blocks | No |
| 82–83 | GID that can use reserved blocks | No |
| 84–87 | First non-reserved inode in file system | No |
| 88–89 | Size of each inode structure | Yes |
| 90–91 | Block group that this superblock is part of (if backup copy) | No |

Superblock: details (3/3)

| | | |
|----------|---|-----|
| 92–95 | Compatible feature flags (see Table 15.6) | No |
| 96–99 | Incompatible feature flags (see Table 15.7) | Yes |
| 100–103 | Read only feature flags (see Table 15.8) | No |
| 104–119 | File system ID | No |
| 120–135 | Volume name | No |
| 136–199 | Path where last mounted on | No |
| 200–203 | Algorithm usage bitmap | No |
| 204–204 | Number of blocks to preallocate for files | No |
| 205–205 | Number of blocks to preallocate for directories | No |
| 206–207 | Unused | No |
| 208–223 | Journal ID | No |
| 224–227 | Journal inode | No |
| 228–231 | Journal device | No |
| 232–235 | Head of orphan inode list | No |
| 236–1023 | Unused | No |

Demo/exercise

- 1 use `dumpe2fs/fsstat` to check where the superblock copies are
- 2 dump the first three, and compare them: are they equal?
 - note: the main SB starts at offset 1024, the others at 0
 - are the 2nd and 3rd equal?

Group descriptors: details

| Byte Range | Description | Essential |
|------------|--|-----------|
| 0–3 | Starting block address of block bitmap | Yes |
| 4–7 | Starting block address of inode bitmap | Yes |
| 8–11 | Starting block address of inode table | Yes |
| 12–13 | Number of unallocated blocks in group | No |
| 14–15 | Number of unallocated inodes in group | No |
| 16–17 | Number of directories in group | No |
| 18–31 | Unused | No |

Outline

- 1 Introduction
- 2 Layout
- 3 Inodes**
- 4 Directories
- 5 Extended attributes

Inodes

Inodes are the **primary metadata structure**

- **fixed size**, defined in the superblock
 - minimum 128 bytes
 - extra-space can be used to store **extended attributes** (otherwise, data blocks are used)
 - small non-user content can be stored inside direct-block pointer array; e.g., symlink values
- each inode has an **“address”**, starting with 1 (at index 0)
- the **inodes of each block-group are stored in a table**
 - whose location is specified in the block descriptor
- **allocation status** determined using the inode **bitmap**
 - whose location is specified in the block descriptor

Demo/exercise

In our example, which is the size of an inode?

Standard commands to inspect the FS

Standard commands can be used to inspect the FS:

- `-i` in `ls`, shows inode-numbers
- `stat` shows some metadata
- `debugfs` is a FS “debugger” that can display *a lot* of information

Demo/Exercise: Hard-links vs Soft/Sym-links

Check the inode-numbers when creating hard vs soft links

Reserved inodes

Inodes 1 to 10 are reserved

- 1 is (was?) used for keeping track of bad blocks
- 2 is used for the root directory
- 8 is typically for the journal

Inode: details (1/2)

| Byte Range | Description | Essential |
|------------|---|-----------|
| 0–1 | File mode (type and permissions) (see Tables 15.11, 15.12, and 15.13) | Yes |
| 2–3 | Lower 16 bits of user ID | No |
| 4–7 | Lower 32 bits of size in bytes | Yes |
| 8–11 | Access Time | No |
| 12–15 | Change Time | No |
| 16–19 | Modification time | No |
| 20–23 | Deletion time | No |
| 24–25 | Lower 16 bits of group ID | No |
| 26–27 | Link count | No |
| 28–31 | Sector count | No |
| 32–35 | Flags (see Table 15.14) | No |
| 36–39 | Unused | No |
| 40–87 | 12 direct block pointers | Yes |
| 88–91 | 1 single indirect block pointer | Yes |
| 92–95 | 1 double indirect block pointer | Yes |
| 96–99 | 1 triple indirect block pointer | Yes |

Inode: details (2/2)

| | | |
|---------|---|----|
| 100–103 | Generation number (NFS) | No |
| 104–107 | Extended attribute block (File ACL) | No |
| 108–111 | Upper 32 bits of size / Directory ACL Yes / | No |
| 112–115 | Block address of fragment | No |
| 116–116 | Fragment index in block | No |
| 117–117 | Fragment size | No |
| 118–119 | Unused | No |
| 120–121 | Upper 16 bits of user ID | No |
| 122–123 | Upper 16 bits of group ID | No |
| 124–127 | Unused | No |

Outline

- 1 Introduction
- 2 Layout
- 3 Inodes
- 4 Directories**
- 5 Extended attributes

Directories


The old directory entry was a simple data structure containing

- the **file name** (variable length, 1-255 chars)
 - the length of an entry is rounded up to a multiple of four
- the **inode number** (AKA index+1)

A)



B)



The newer one uses **one byte in the filename-length to store the file type**. This could be used to detect when an inode has been reallocated

Directories (new version): details

| Byte Range | Description | Essential |
|------------|-----------------------------|-----------|
| 0–3 | Inode value | Yes |
| 4–5 | Length of this entry | Yes |
| 6–6 | Name length | Yes |
| 7–7 | File type (see Table 15.24) | No |
| 8+ | Name in ASCII | Yes |


file type:

- 1 Regular file
- 2 Directory
- 3 Character device
- 4 Block device
- 5 FIFO
- 6 Unix Socket
- 7 Symbolic link

Hash trees

In ext3 and ext4, option `dir_index` allows the FS to use hashed B-trees to speed up name lookups

- similar to NTFS
- a ro-compatible feature flag will be set in the superblock
- still the same directory-entry structures, but in sorted order
 - node descriptors are “hidden” like deleted files



Exercises (1/2)

Demo/exercise

If your root FS is ext4, verify the inode number of the root directory of your system

Demo/exercise

List files in the root directory of `linus-ext2.dd` by

- (ro) mounting it
- using TSK's `fls`
- using `ImHex`
 - get the block-size from the SB
 - the GD-table is in the block after the SB
 - decode the first GD and get the block for the (first) inode-table
 - the inode `n.2` (index 1) is for the root directory
 - the first direct-block-pointer points to the beginning of the directory
 - decode the directory entries

Demo/exercise

List the data block for `/pics/linus/linus_2018.jpg`, by using

① TSK

- output can be confusing: “Direct Blocks” are all blocks, while “Indirect Blocks” are the blocks containing the indirect pointers and so on

② ImHex; *beware*: the inode number for `/pics` is 5033

i.e., it's the first inode of Group 18 (index 17); since, $17 * 296 = 5032$

Outline

- 1 Introduction
- 2 Layout
- 3 Inodes
- 4 Directories
- 5 Extended attributes**

Extended attributes

- **extended attributes** are a list of **name and value pairs**
 - names are in the form *namespace.attribute-name* form
 - current namespaces are security, system, trusted, and **user**
 - one use is implementing POSIX ACLs, see **setfacl**
 - kernel and filesystem may place limits on the maximum number and size of extended attributes that can be associated with a file
- to get/set attributes: commands **getfattr**/**setfattr**
- the attributes are stored in a block allocated to the file
 - If more files have the same attributes, they can share the same block

See `xattr(7)`

Extended attributes: details




Table 15.15. Data structure for the extended attribute header.

| Byte Range | Description | Essential |
|------------|------------------------|-----------|
| 0–3 | Signature (0xEA020000) | No |
| 4–7 | Reference count | No |
| 8–11 | Number of blocks | Yes |
| 12–15 | Hash | No |
| 16–31 | Reserved | No |

Table 15.16. Data structure for the extended attribute name entries.

| Byte Range | Description | Essential |
|------------|----------------------------------|-----------|
| 0–0 | Length of name | Yes |
| 1–1 | Attribute type (see Table 15.17) | Yes |
| 2–3 | Offset to value | Yes |
| 4–7 | Block location of value | Yes |
| 8–11 | Size of value | Yes |
| 12–15 | Hash of value | No |
| 16+ | Name in ASCII | Yes |

Extended attributes

- List the extended attributes by ro-mounting the image and `getfattr`
- Try to decode them using `ImHex`
 - hint: one is “easy”, but others are tricky; why?

- [Car05] Brian Carrier.
File System Forensic Analysis.
Addison-Wesley Professional, 2005.
- [Fai12] Kevin D. Fairbanks.
An analysis of ext4 for digital forensics.
Digital Investigation, 9:S118–S130, 2012.
The Proceedings of the Twelfth Annual DFRWS Conference.